

Software Defined Networking

MOBILE AND CYBER-PHYSICAL SYSTEMS

FEDERICA PAGANELLI

federica.paganelli@unipi.it

Learning Objectives

- Requirements and current network technology limitations
- Software Defined Networking (SDN)
 - Main concepts and definitions
 - SDN Data Plane
 - SDN Control Plane
 - Some Applications of SDN



Requirements and current network technology limitations

Evolving Network Requirements

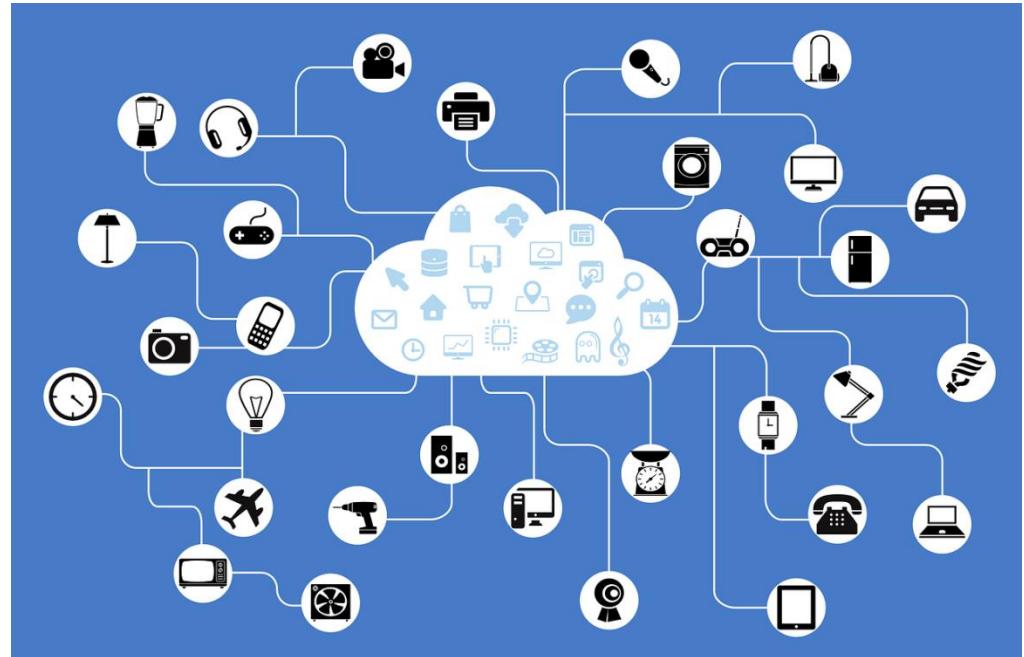
Why redesigning traditional approaches to network architecture?

Demand is increasing

- Cloud computing
- Big Data
- Mobile traffic
- Internet of Things (IoT)

... just to mention main trends demanding for network architecture redesign

From distributed
to centralized
logical approach



⇒ not just matter of ~~the netw~~ technologies

Evolving Network Requirements

But... also supply is increasing

the capacity of network technologies is capable of absorbing rising loads

performance of network devices has increased

- Faster processor speed, greater buffer capacity and buffer access speed (thanks to larger, faster memories)

So what's the problem?

It is not just a matter of supply and demand

- ① traffic patterns have changed and become more dynamic and complex

problem } smaller volumes
- ~1 pattern changes

Traffic patterns measure how traffic varies in time and space

Complex and dynamic traffic patterns 1/3

②

Modern distributed applications typically access **multiple databases** and **servers** that must communicate with each other NOT ONLY

MULTIPLE REQUESTS



“horizontal” traffic between servers in addition to the “vertical” client/server one

③

Unified communications (UC) services are increasingly used within enterprises.

↳ Attende using portal x integrated services per user
in 1 platform
↳

- integration of communication services such as instant messaging (chat), presence information, voice, mobility, audio, web and video traffic into a unique service delivery platform
- applications have thus to access multiple servers

Cooperation
with
several
servers

Complex and dynamic traffic patterns 2/3

4

Widespread use of clouds

setup net with local db but use external cloud
↑
↓
LOT
TRAFFIC
IN THE

- a significant amount of traffic that was previously local to enterprise networks has moved to remote clouds
- creating increased and often very **unpredictable loads on enterprise routers.**

ENTER NET

Q

Cloud ↔ ENTER.
NET

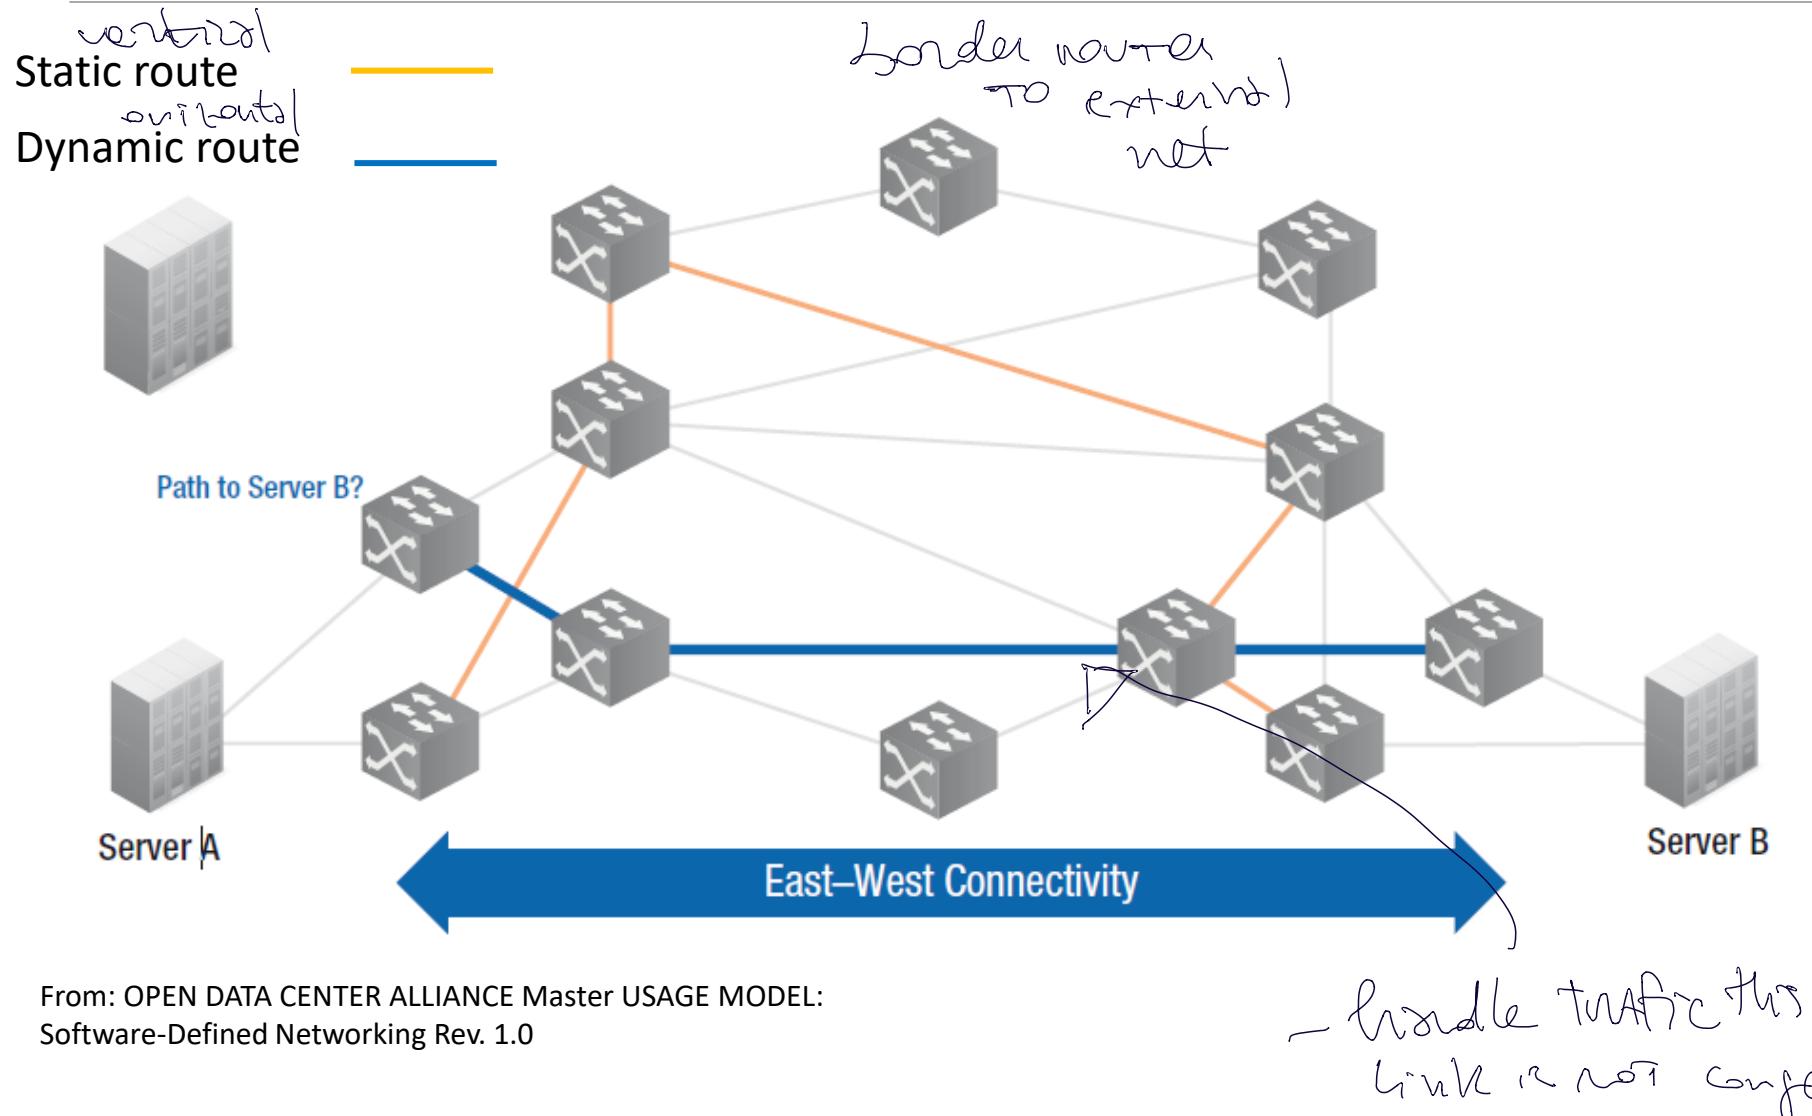
5

Server virtualization.

It allows partitioning a single hardware machine into multiple virtual ones (e.g. Virtual Machines).

- considerable amount of **traffic among virtual servers** (e.g., **migration**, switchover). ↗
- such **horizontal traffic flow changes in location and intensity over time**, demanding a **flexible approach to managing network resources.**

Traffic between servers



By some estimates, **over 70 percent of traffic is now East-West** (traffic between servers within the data center).

The traditional three-layer data center network design (core-aggregation-access) is ill-equipped to handle this kind of traffic

Complex and dynamic traffic patterns 3/3

mobility support

The increasing use of mobile devices anywhere and any time

- fast-changing and unpredictable large traffic loads on the network
- the network attachment point can rapidly change

web navigation (no latency sensitive)
sensitive to latency (medical) IOT app. (Zigbee QoS)

Different QoS requirements demanded by a variety of

applications \Rightarrow detect if flow is HTTP or other
and apply different policy

- differentiated treatment of granular traffic flows

to handle traffic
just based
on ~~MAC~~
destination
address

Limitations of current networks

Why is change needed?
→ managing net res.
more flexibility
(routers, switches,
flows)

All these emerging trends require a **flexible management of network resources**

... but traditional networks show some **limitations**...

protocols
for routing, control,
switching

- ① **Static, complex architecture:** It consists in a number of independently defined protocols each of which addresses a portion of networking requirements -> **complexity that is difficult to master!**
- ② **Inconsistent policies:** manual procedures needed to change configuration in multiple switches, routers, firewalls -> error prone procedures to routers & firewalls manual procedures
- ③ **Vendor dependence:** enterprises and carriers need to deploy new capabilities and services rapidly in response to changing business needs and user demands. A lack of open interfaces for network functions leaves the enterprises limited by the relatively slow product cycles of vendor equipment
lack of interfaces

→ SDN/NFV
from: Open Networking Foundation (ONF)
but miss interoperability of vendors ⇒ IT's own features

Key to Internet Success: Layers

we're dealing with
plain text

Applications

...built on...

Reliable (or unreliable) transport

...built on...

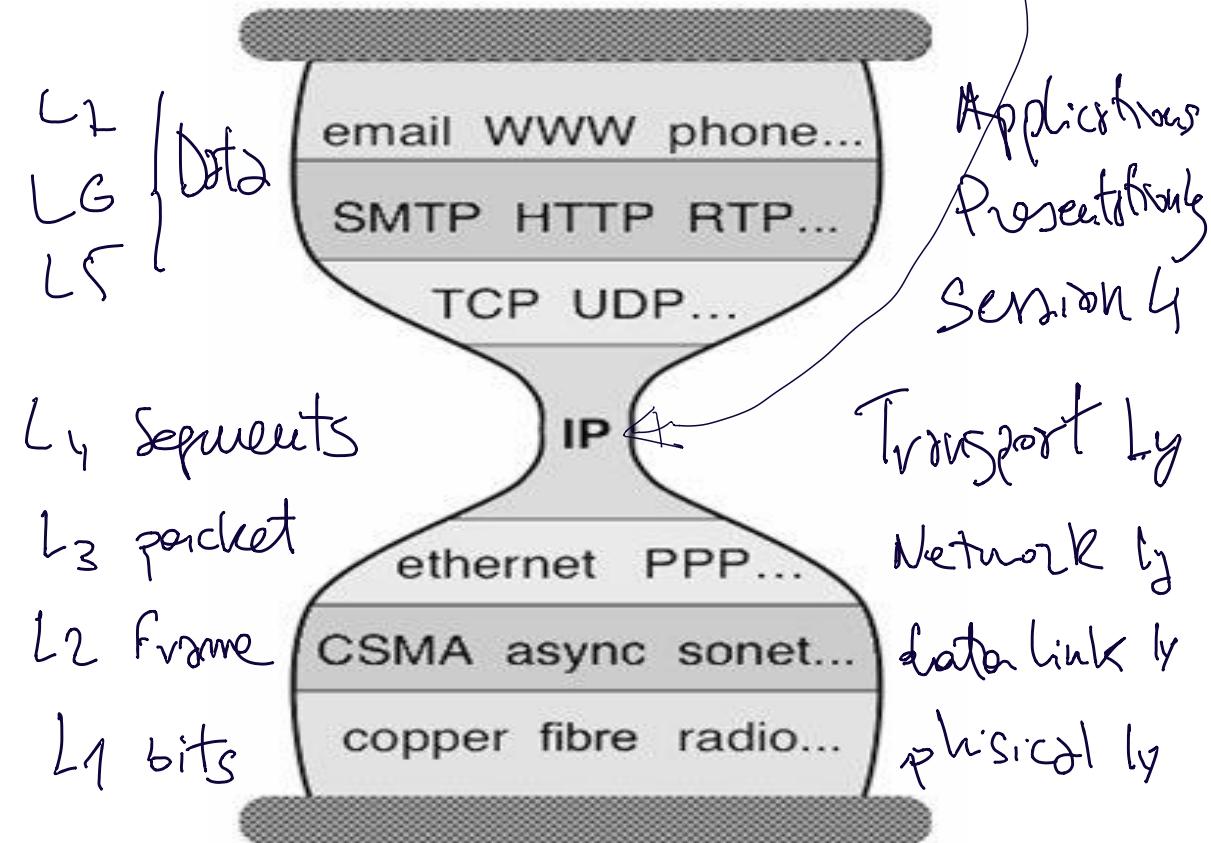
Best-effort global packet delivery

...built on...

Best-effort local packet delivery

...built on...

Physical transfer of bits



Why Is Layering So Important?

Decomposed delivery into fundamental components



Independent but compatible innovation at each layer

A practical success of unprecedented proportions...

internet

in not seen as date plane

Network Control/Management Plane

→ when problem arise
⇒ new protocol
(errors => ICMP)

no general principles or abstractions
guiding the design of network
control/management plane

- in contrast, e.g., to distributed systems or database systems



Control plane is basically composed of various distributed “control” protocols

- For each new feature/functionality needed, a new protocol is designed

Network boxes are “closed” equipment

- (mostly proprietary) software bundled with hardware
- vendor-specific interfaces
 - ⇒ very difficult to manage
- Big bag of protocols
- Notoriously difficult to manage
- Slow protocol standardization
- Slow innovation!

Software Defined Networking

MAIN CONCEPTS AND DEFINITIONS

Network-layer functions

Goal to define principles
⇒ address req. in extending
- traffic
Flowst
very granular
Fast timescales level
(per-packet)

- **forwarding:** move packets from router's input to appropriate router output RONT segment part of source / dest **data plane**
 ↳ Decide at port according to Forward Table
- **routing:** determine route taken by packets from source to destination Compute forwardly TABLE **control plane**

Slow time-scales
(per control event)

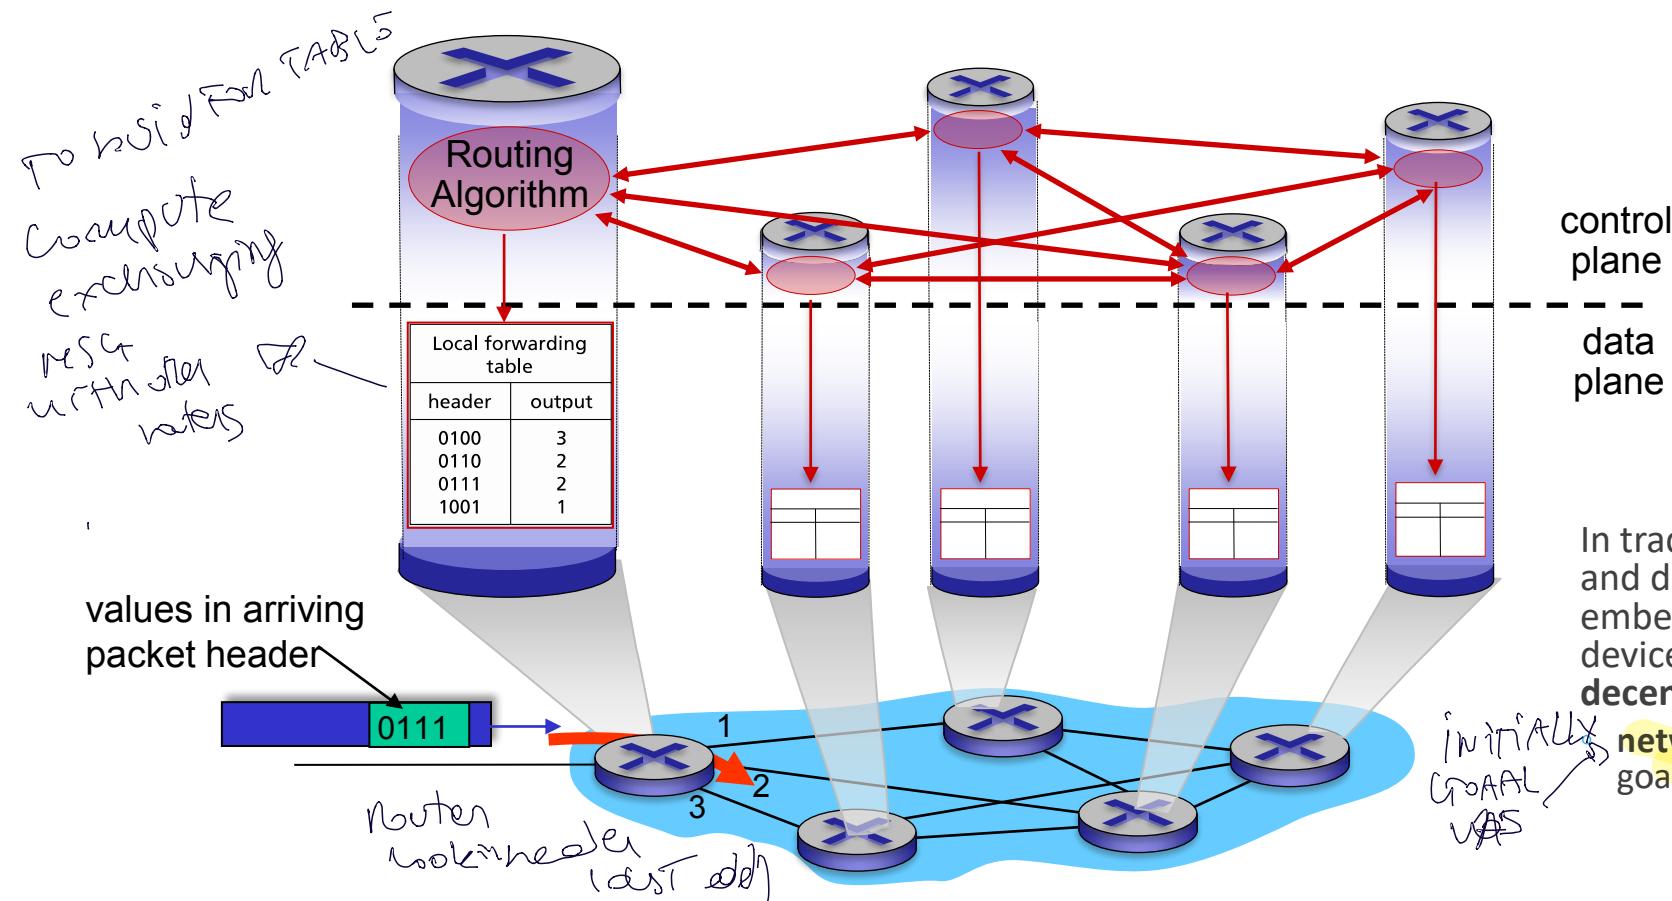
Two approaches to structuring network control plane:

i) **per-router control (traditional)** DISTRIBUTED

ii) **logically centralized control (software defined networking)** typically centralized implementation of control

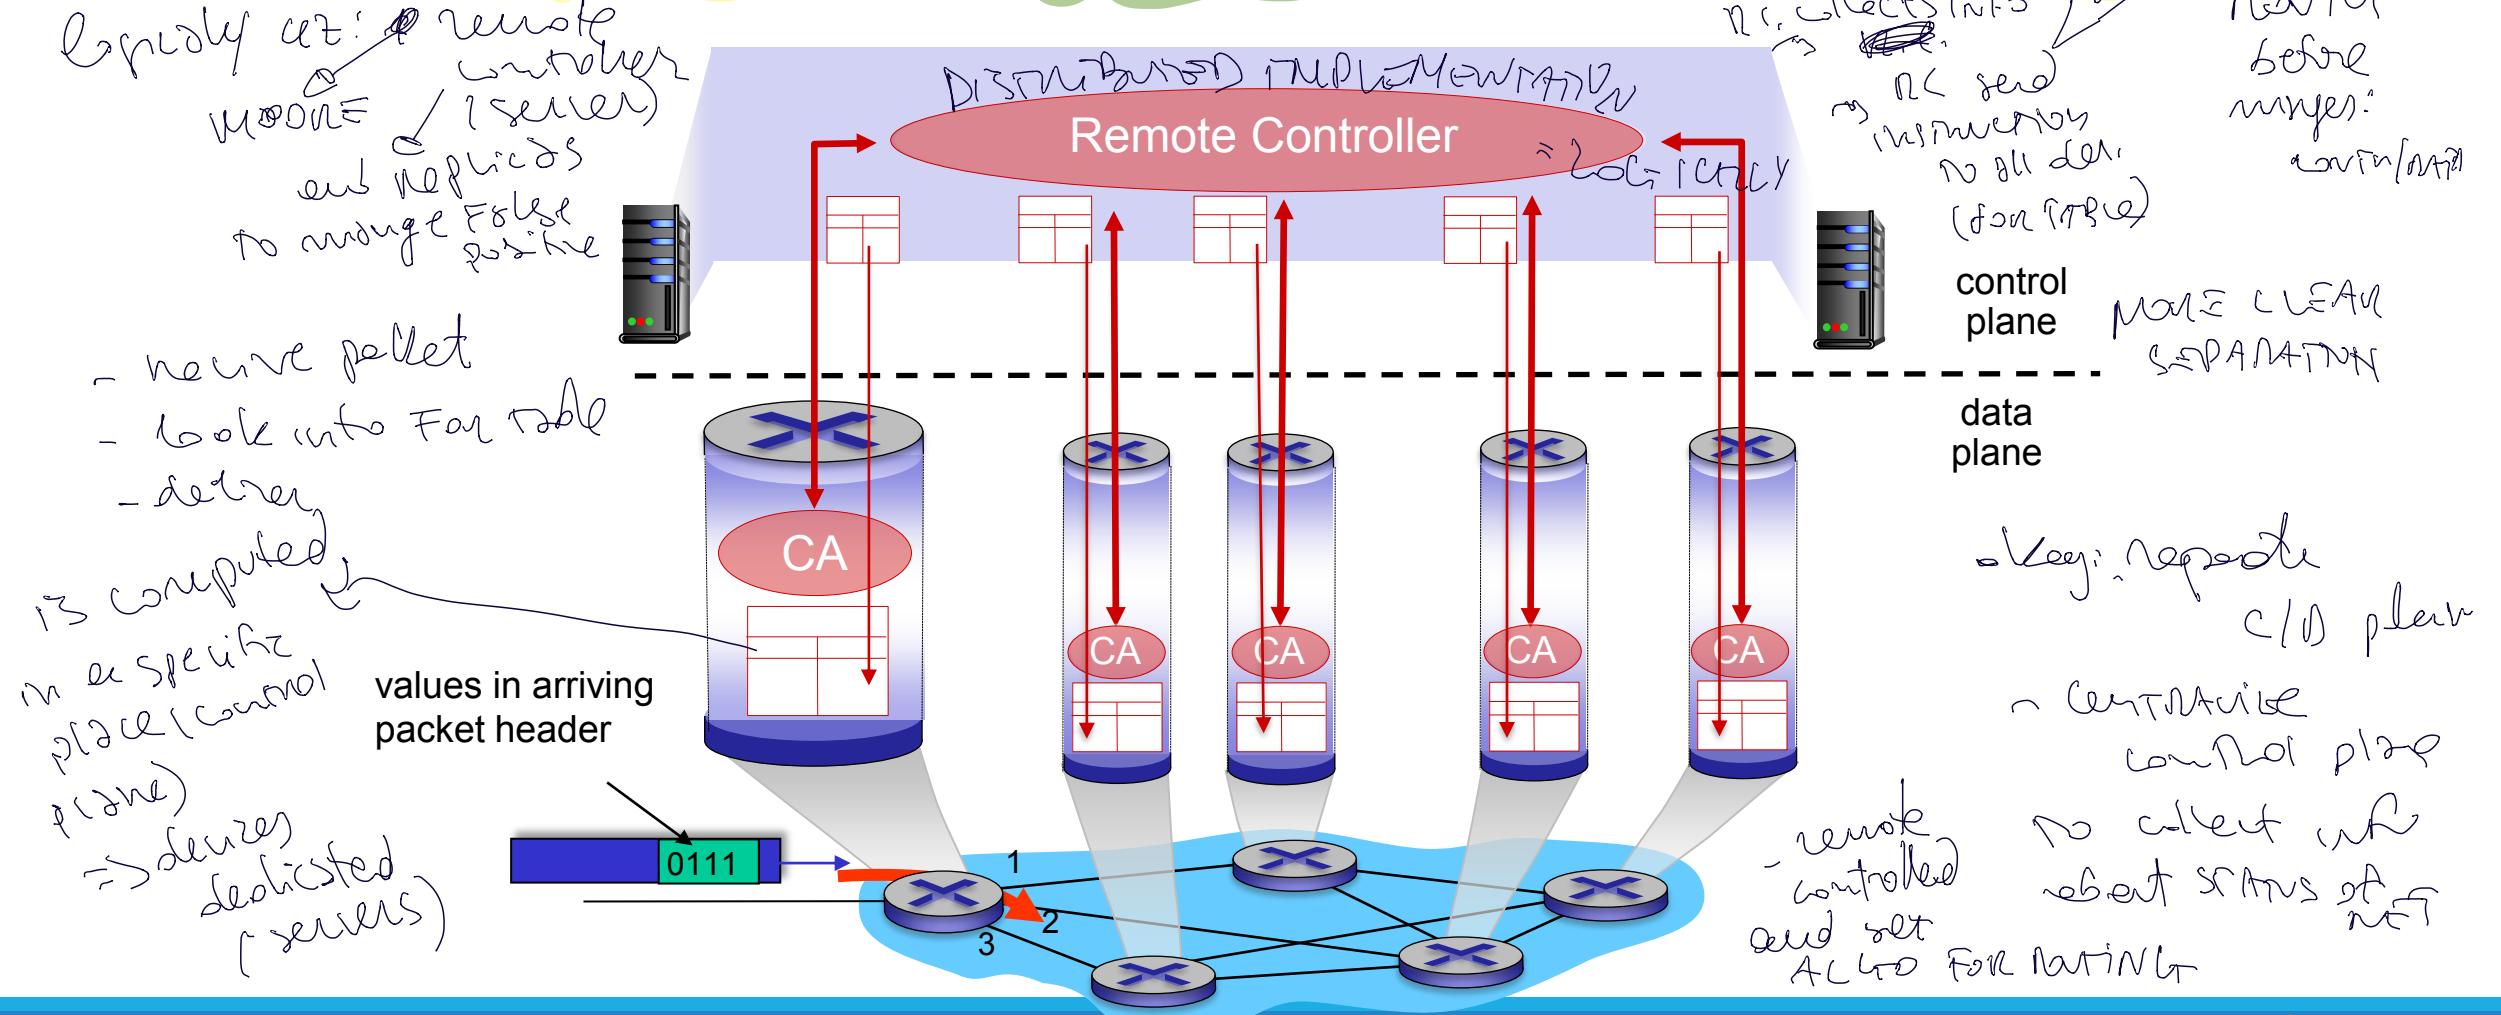
Per-router control plane

Individual routing algorithm components *in each and every router* interact in the control plane to compute forwarding tables



Software-Defined Networking (SDN) control plane

Remote controller computes, installs forwarding tables in routers



③ - manage + strategy
of why

Software defined networking (SDN)

Why a *logically centralized* control plane?

- ① easier network management: avoid router misconfigurations, greater flexibility of traffic flows

Central concept of table-based forwarding (recall OpenFlow API) allows “programming” routers

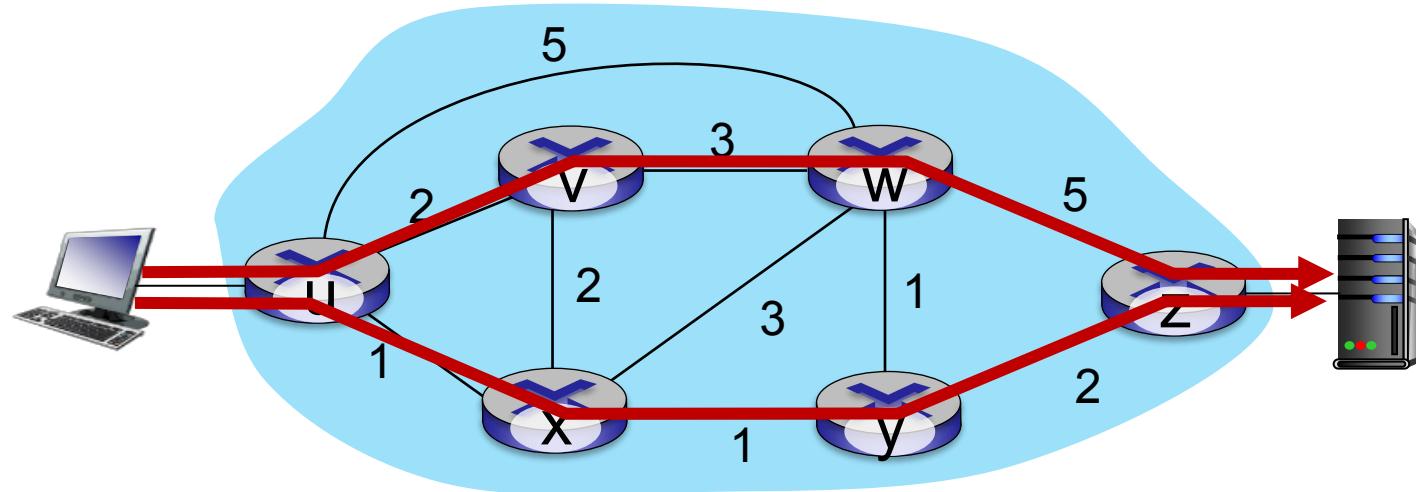
- centralized “programming” easier: compute tables centrally and distribute
- distributed “programming” more difficult: compute tables as result of distributed algorithm (protocol) implemented in each-and-every router

open (non-proprietary) implementation of control plane

- foster innovation: let 1000 flowers bloom

classifying
unit of
FOR TM
⇒ change
behavior
of routers

Traffic engineering: difficult with traditional routing

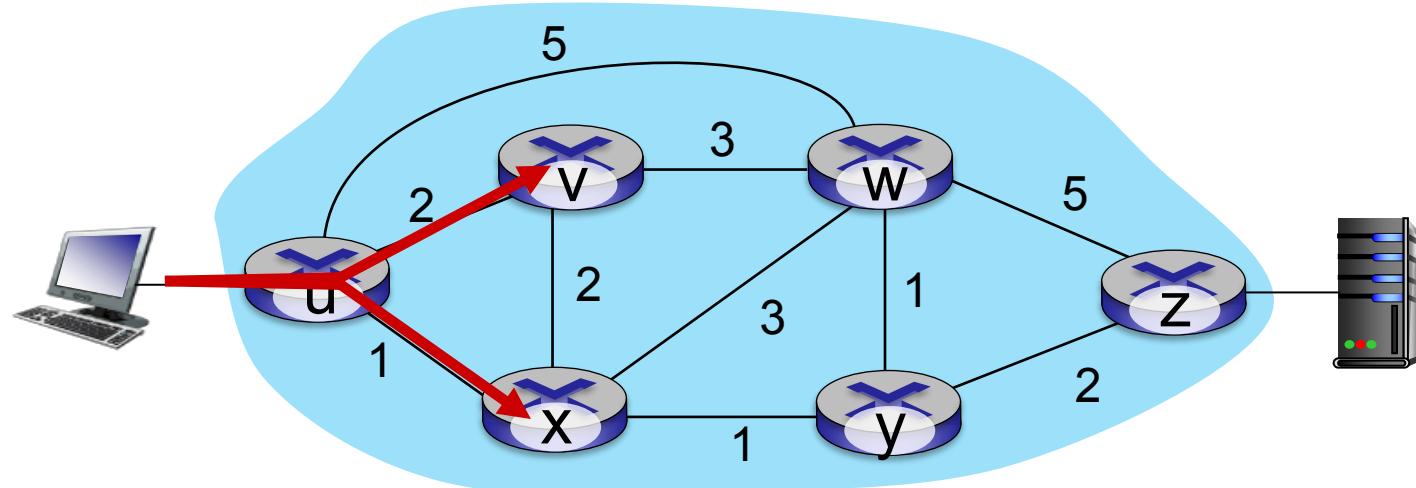


Q: what if network operator wants u-to-z traffic to flow along $uvwz$, rather than $uxyz$?

A: need to re-define link weights so traffic routing algorithm computes routes accordingly (or need a new routing algorithm)!

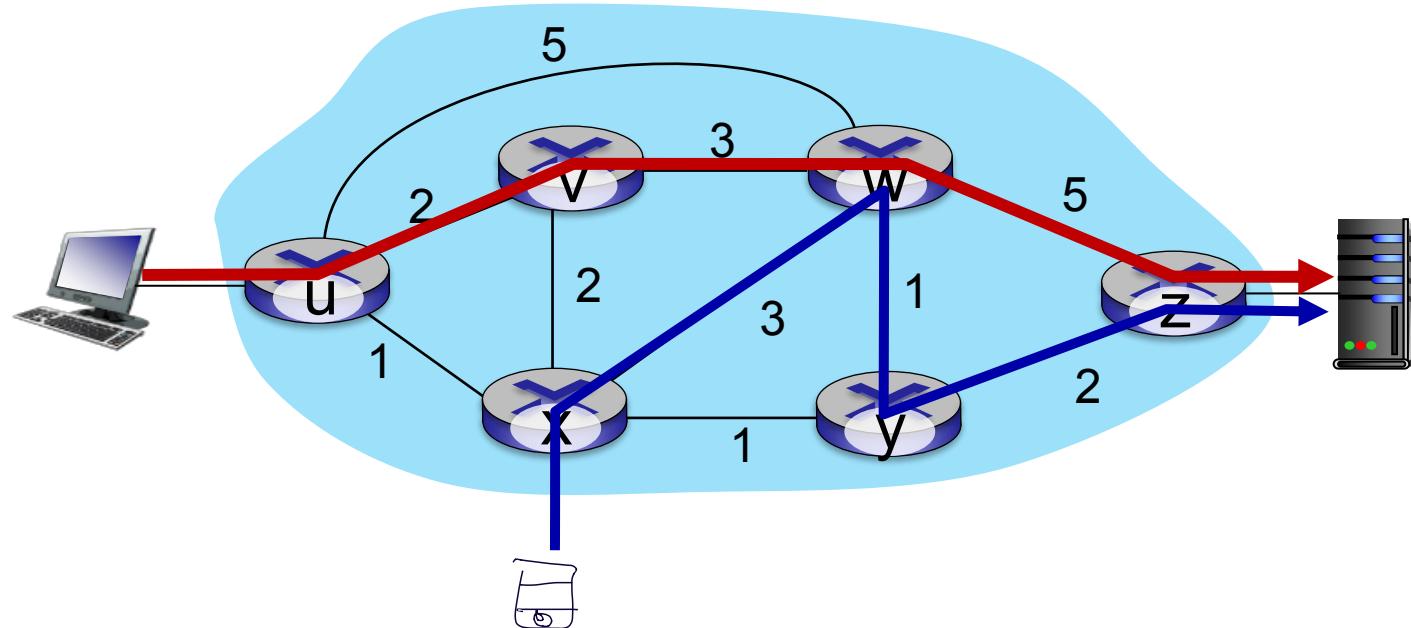
link weights are only control “knobs”: not much control!

Traffic engineering: difficult with traditional routing



Q: what if network operator wants to split u-to-z traffic along uvwz *and* uxyz (load balancing)?
A: can't do it (or need a new routing algorithm)

Traffic engineering: difficult with traditional routing



Q: what if w wants to route blue and red traffic differently from w to z?

A: can't do it (with destination-based forwarding, and LS, DV routing)

We will see how that generalized forwarding and SDN can be used to achieve *any* routing desired

CA: control agent in switches

Software defined networking (SDN)

4. programmable
control
applications

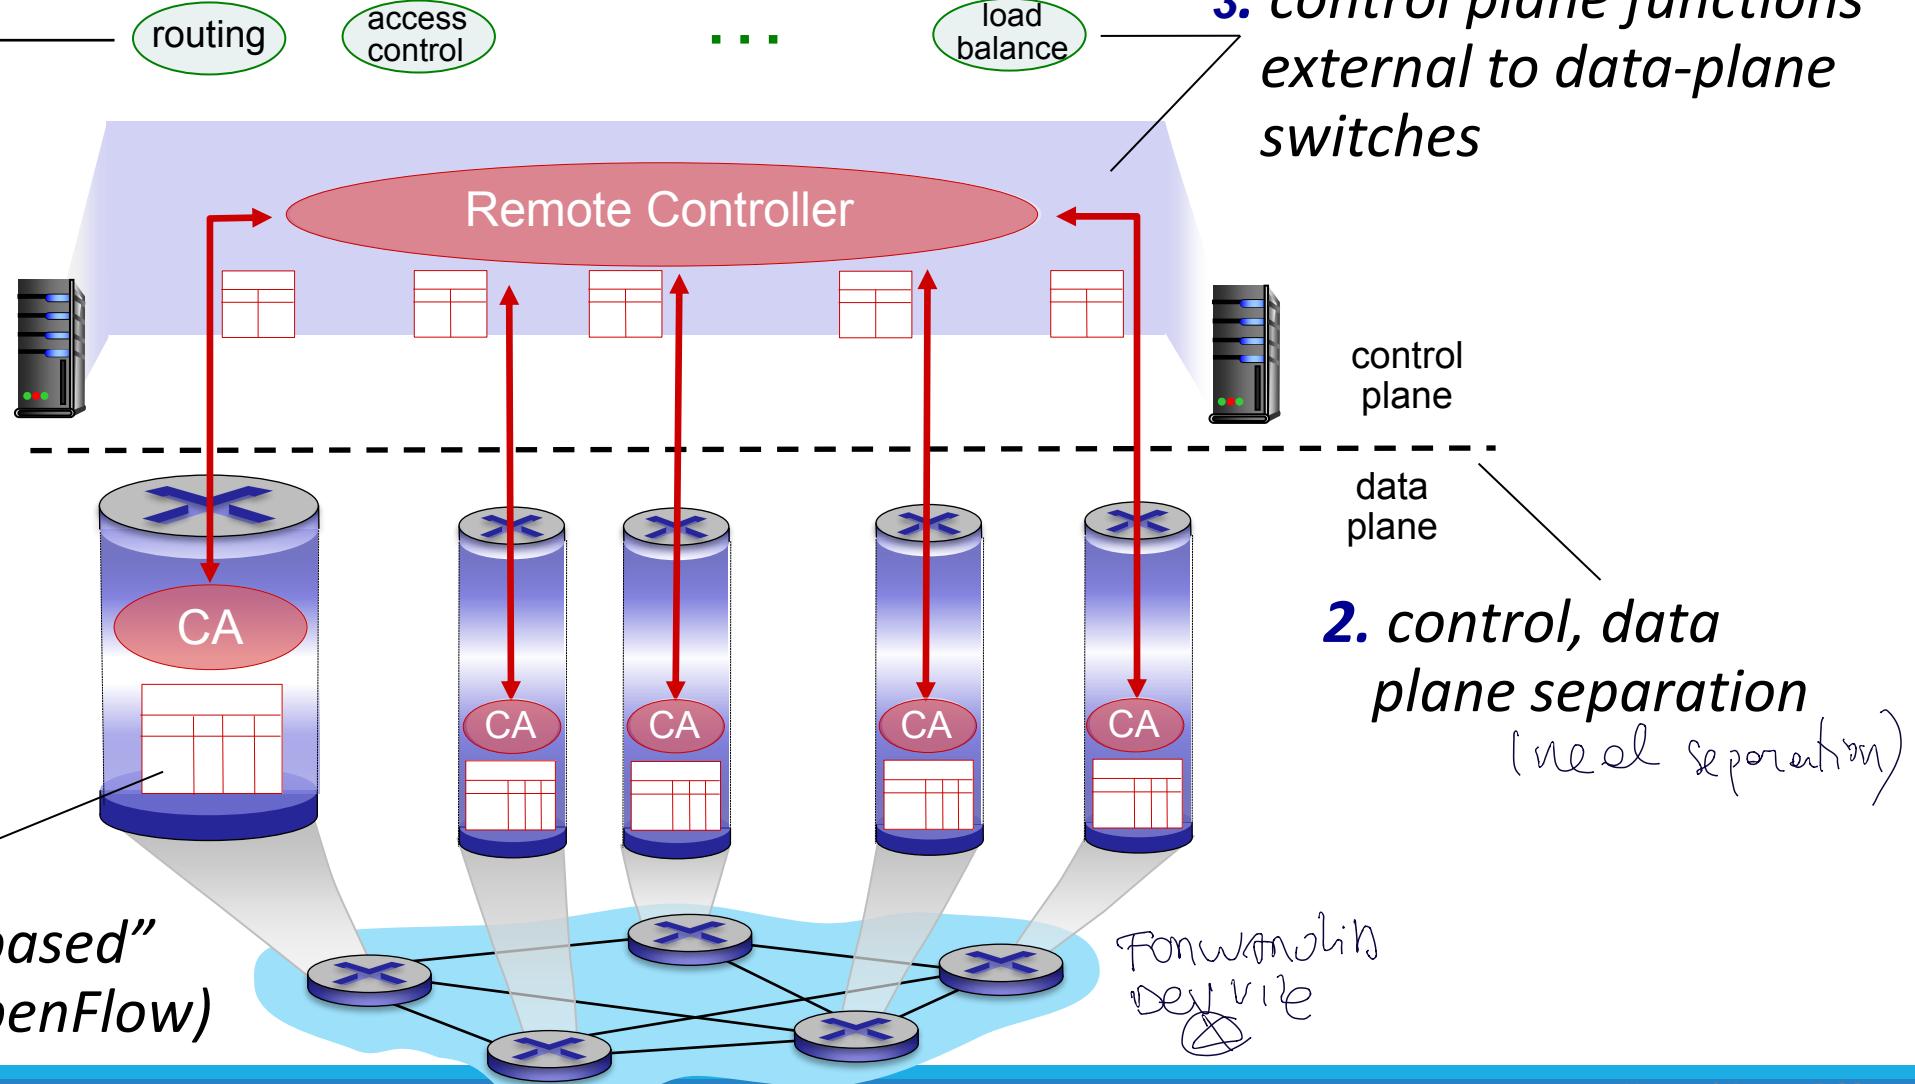
routing

access
control

...

load
balance

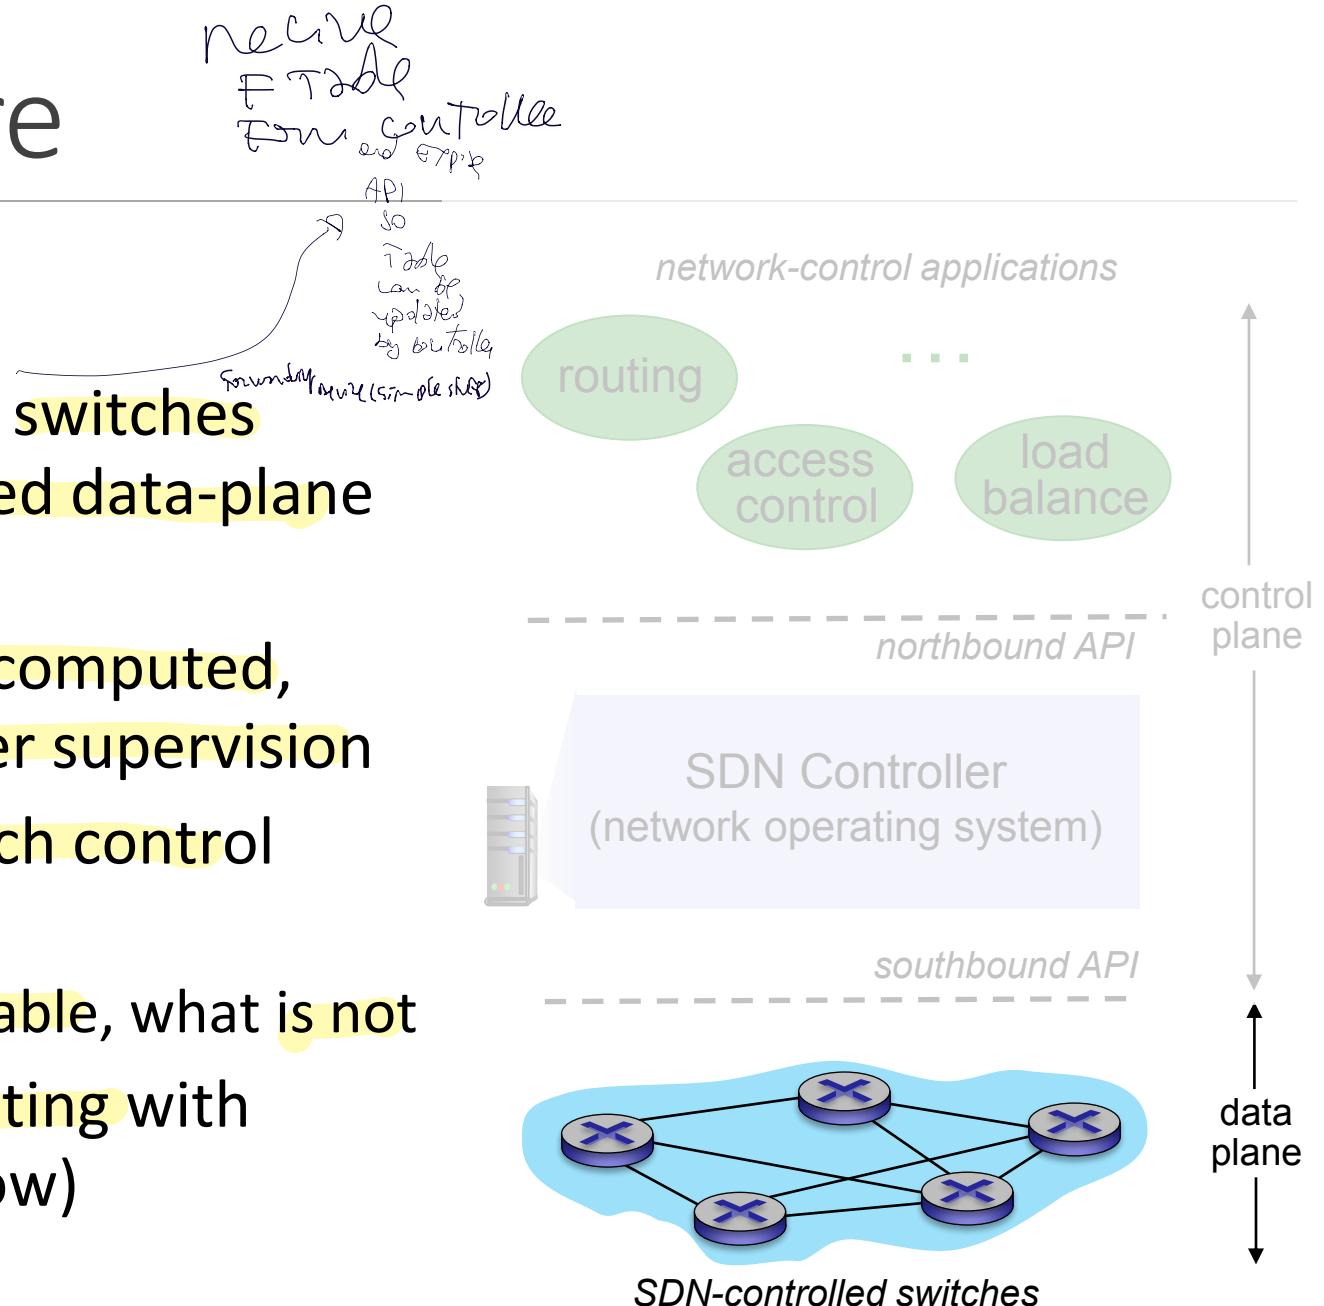
3. control plane functions
external to data-plane
switches



SDN Architecture

Data-plane switches:

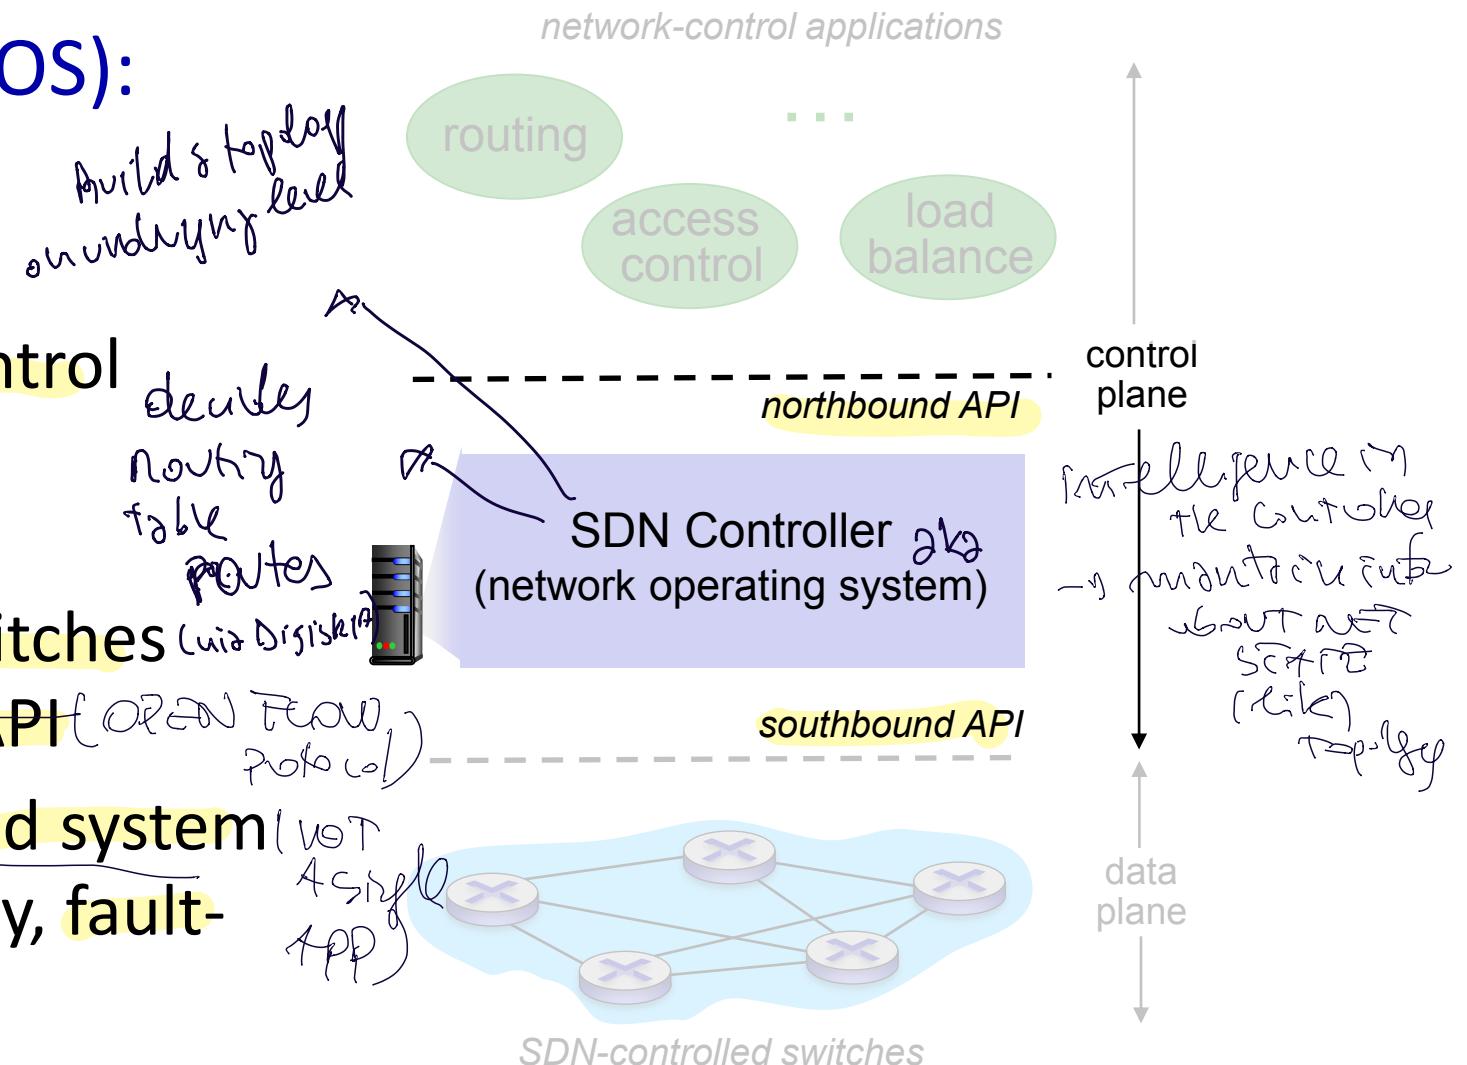
- fast, simple, commodity switches implementing generalized data-plane forwarding in hardware
- flow (forwarding) table computed, installed under controller supervision
- API for table-based switch control (e.g., OpenFlow)
 - defines what is controllable, what is not
- protocol for communicating with controller (e.g., OpenFlow)



Software defined networking (SDN)

SDN controller (network OS):

- maintain network state information
- interacts with network control applications “above” via northbound API
- interacts with network switches “below” via southbound API
- implemented as distributed system for performance, scalability, fault-tolerance, robustness



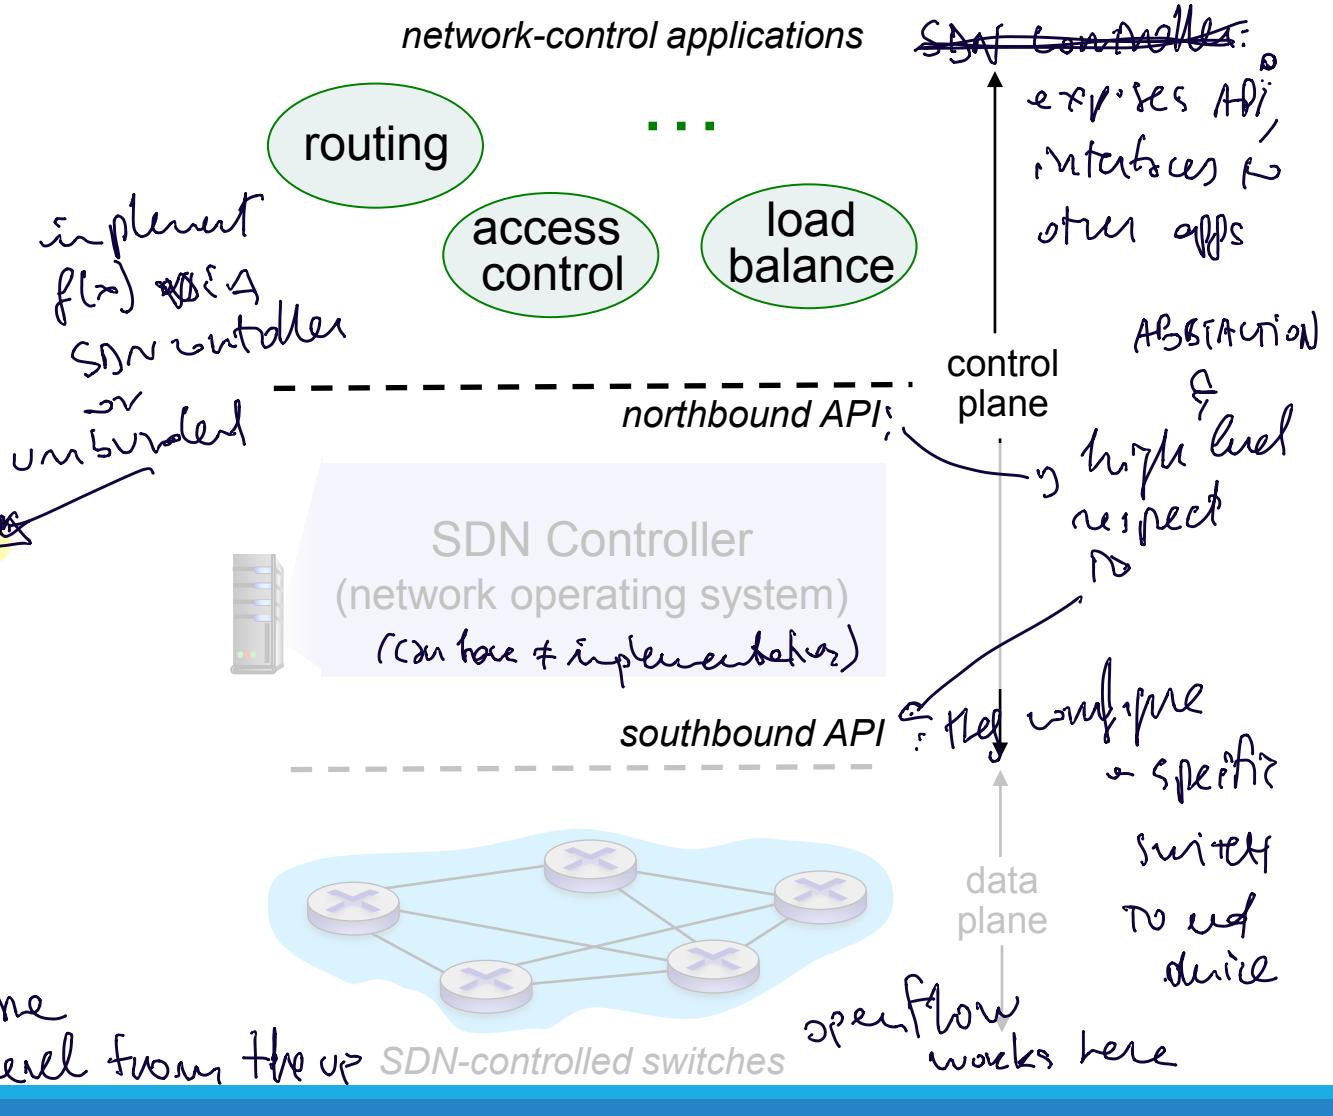
Software defined networking (SDN)

network-control apps:

- “brains” of control:
implement control functions
using lower-level services, API
provided by SDN controller
- unbundled: can be provided by
^(extended) 3rd party: distinct from routing
vendor, on ^{part of} SDN controller

external dev. to operate \Rightarrow
Vendor few (Cisco)
via API

\Rightarrow can specify configuration on the
underlying level from the up ^{SDN-controlled switches} openflow works here



Brief history

~ 2004 Research on new management paradigms

- Princeton, Stanford/Berkeley

2008 Software defined networking

- NOX Network Operating System [NICIRA]
- OpenFlow switch interface [Stanford/NICIRA]

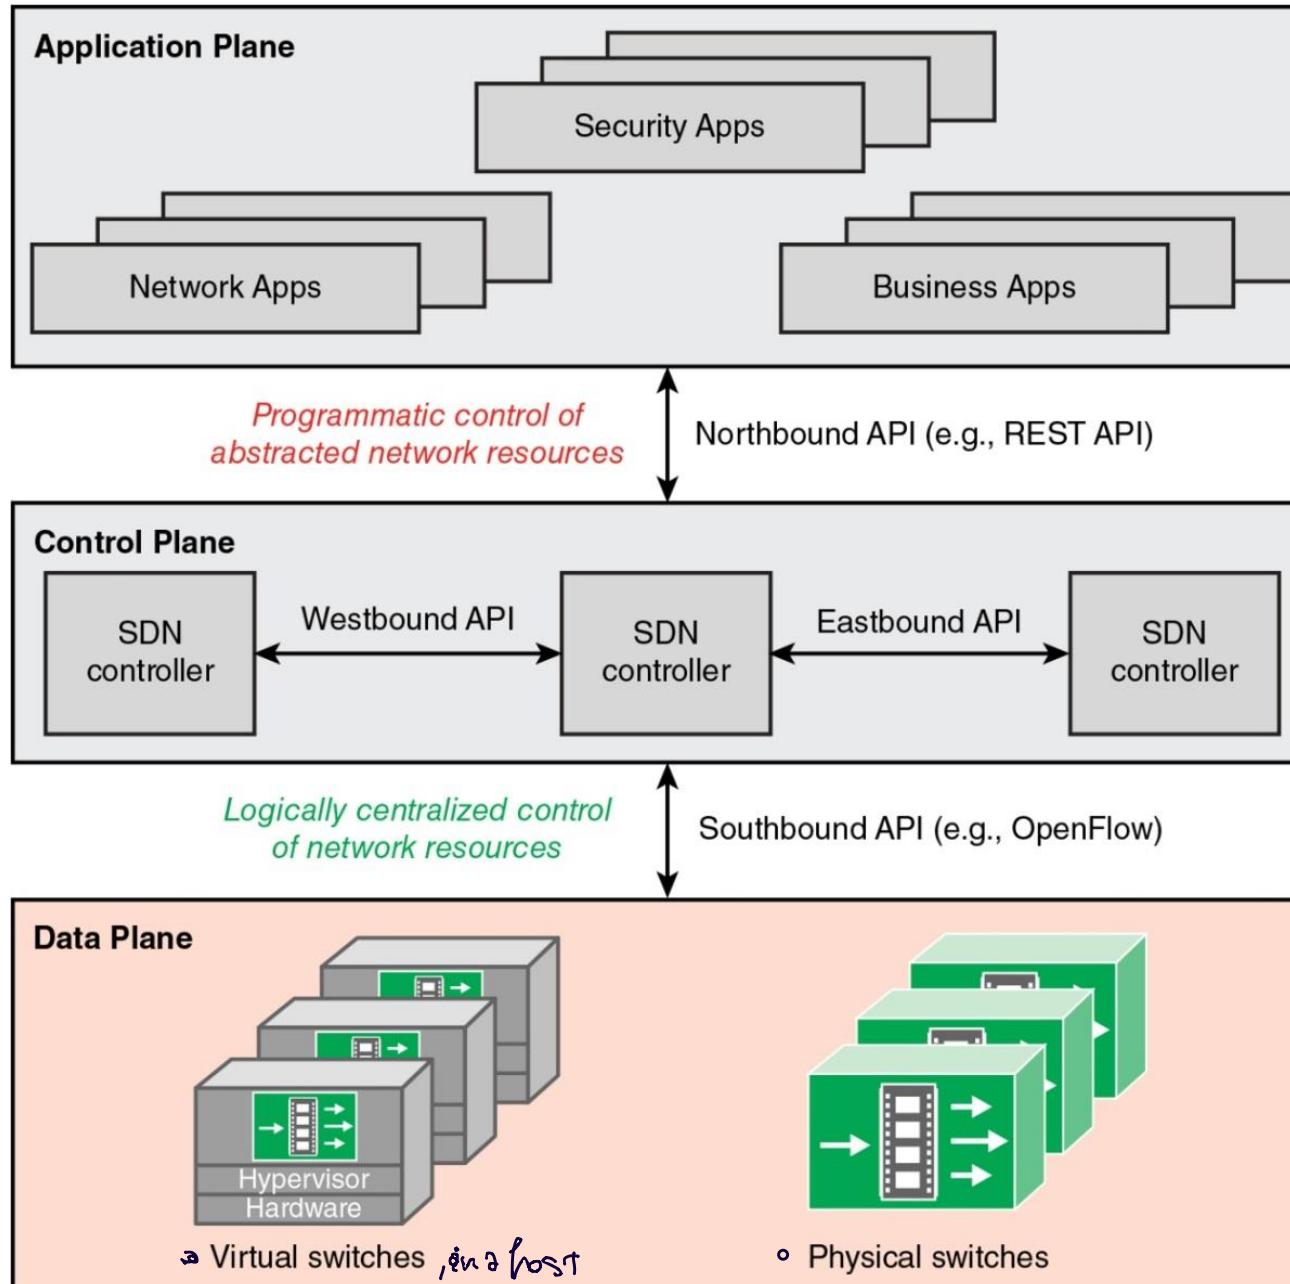
2011 Open Networking Foundation

- Board: Google, Yahoo, Verizon, ecc.
- Members: Cisco, HP, Dell,..

2013

- 1600 attendees at the Open Networking Summit
- SDN used in Google WAN

SDN Data Plane



Data Plane

Resource or infrastructure layer made by forwarding devices

- perform the transport and processing of data according to decisions made by the SDN control plane.
- without embedded software to make autonomous decisions.

-> **Forwarding Abstraction**

Data Plane Network Device

~~PLANE~~ Data ~~V~~forwarding function:

look in the forward table to modify the traffic flow

- Accepts incoming data flows from other network devices and end systems
- forwards them along the data forwarding paths computed and established according to the rules defined by the SDN applications

Control support function

- Interaction with the SDN controller for management of forwarding rules
- Reference standard specifications: **OpenFlow switch protocol.**

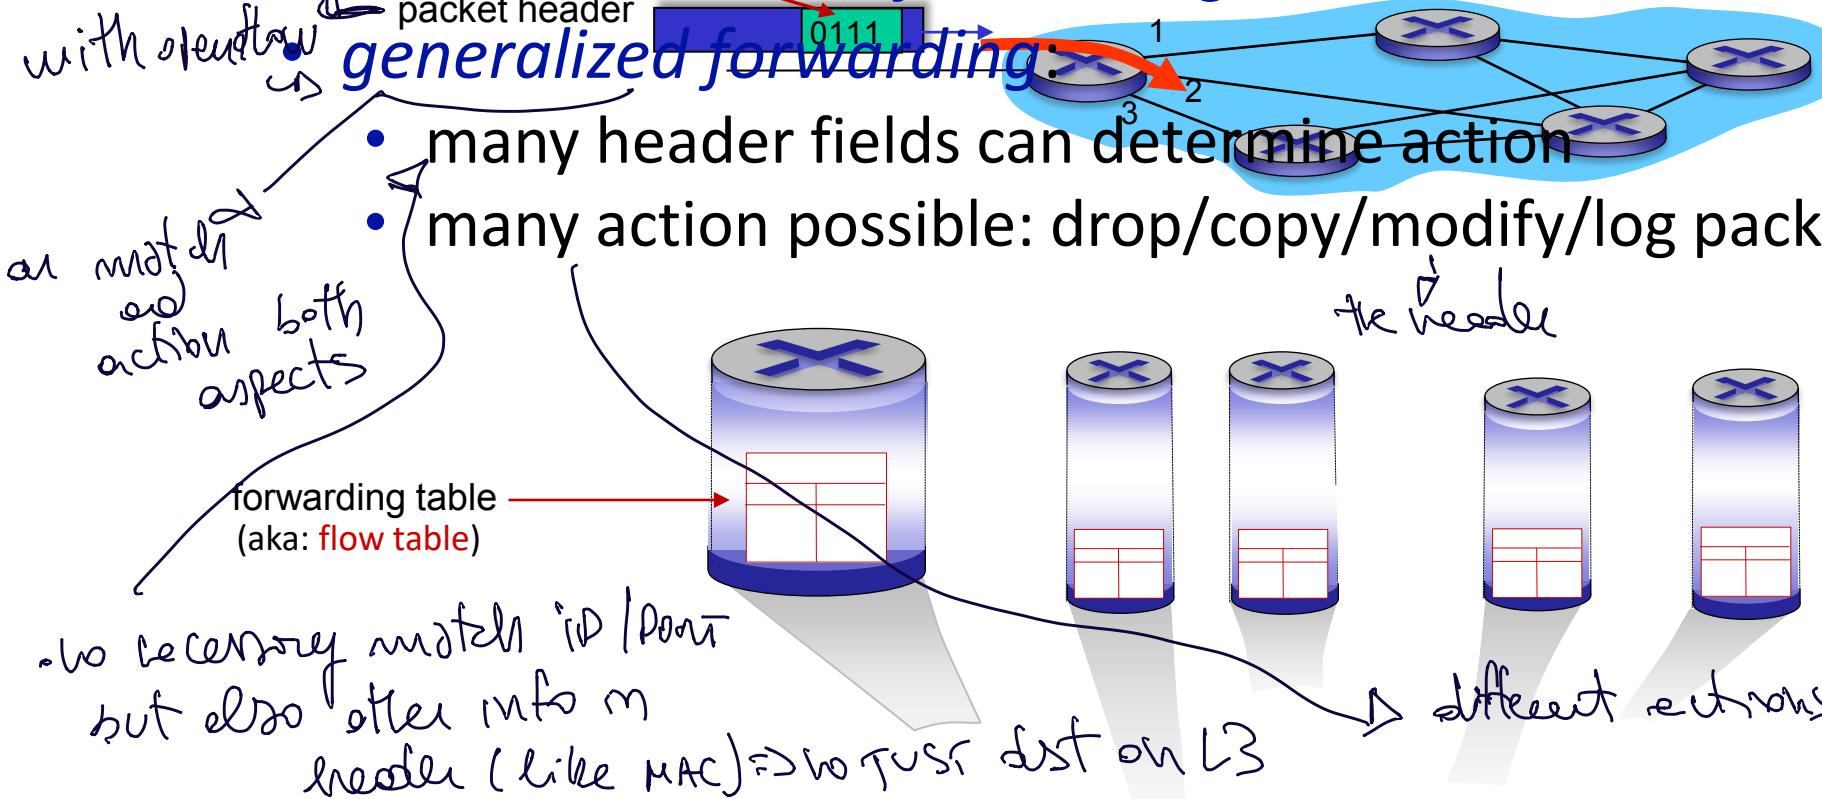
device of data plane must communicate with SDN controller
=> need a module to speak between them

Generalized forwarding: match plus action

Review: each router contains a **forwarding table** (aka: **flow table**)

- “**match plus action**” abstraction: match bits in arriving packet, take action

- **destination-based forwarding**: forward based on dest. IP address



Flow table abstraction

flow: defined by header field values (in link-, network-, transport-layer fields)

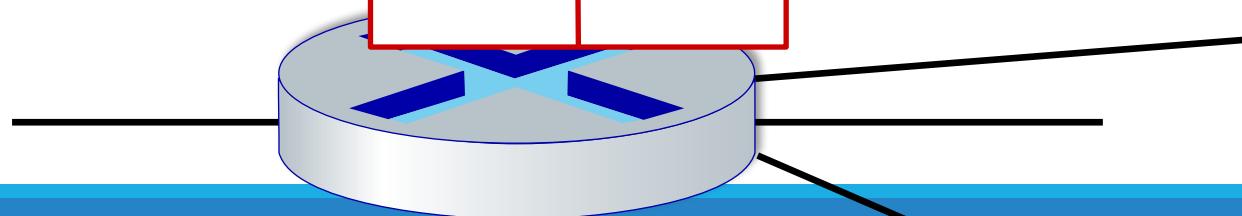
generalized forwarding: simple packet-handling rules

- **match:** patterns values for packet header fields
 - **actions:** for matched packet: drop, forward, modify, matched packet or send matched packet to controller (on packet matching that pattern)
 - **priority:** disambiguate overlapping patterns
 - **counters:** #bytes and #packets (for statistics)
needed by the controller
to take decisions about traffic flow
- stells actions to be performed
on first packet
- can have multiple matches
= SWMS
high priority
- = SWNS
higher priority

Flow table	
match	action

Router's flow table define
router's match+action rules

- Should keep it short as possible



Flow table abstraction

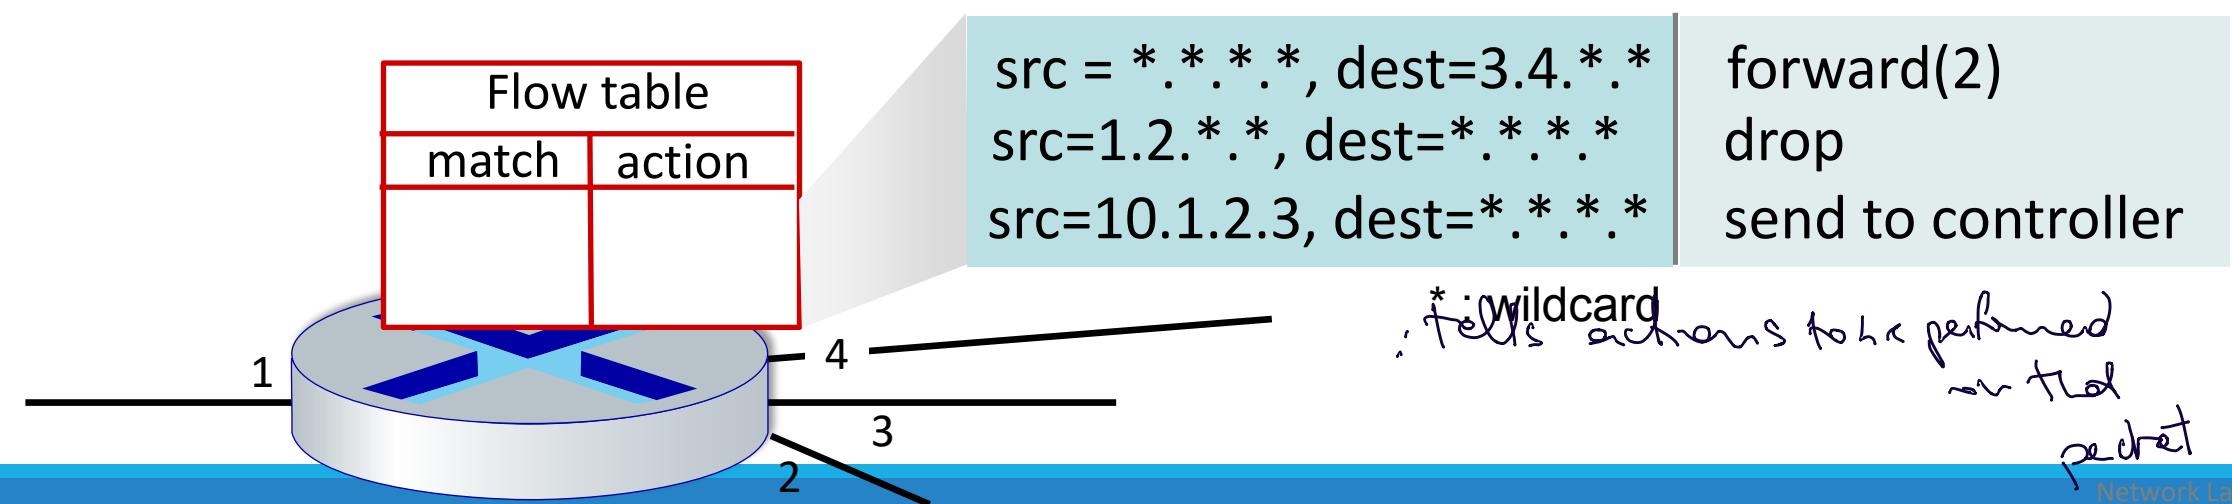
flow: defined by header fields 3

generalized forwarding: simple packet-handling rules

- match: pattern values in packet header fields
- actions: for matched packet: drop, forward, modify, matched packet or send matched packet to controller
- priority: disambiguate overlapping patterns
- counters: #bytes and #packets

if you don't have a specific behavior as switches send
the host will decide actions on this packet

send to controller

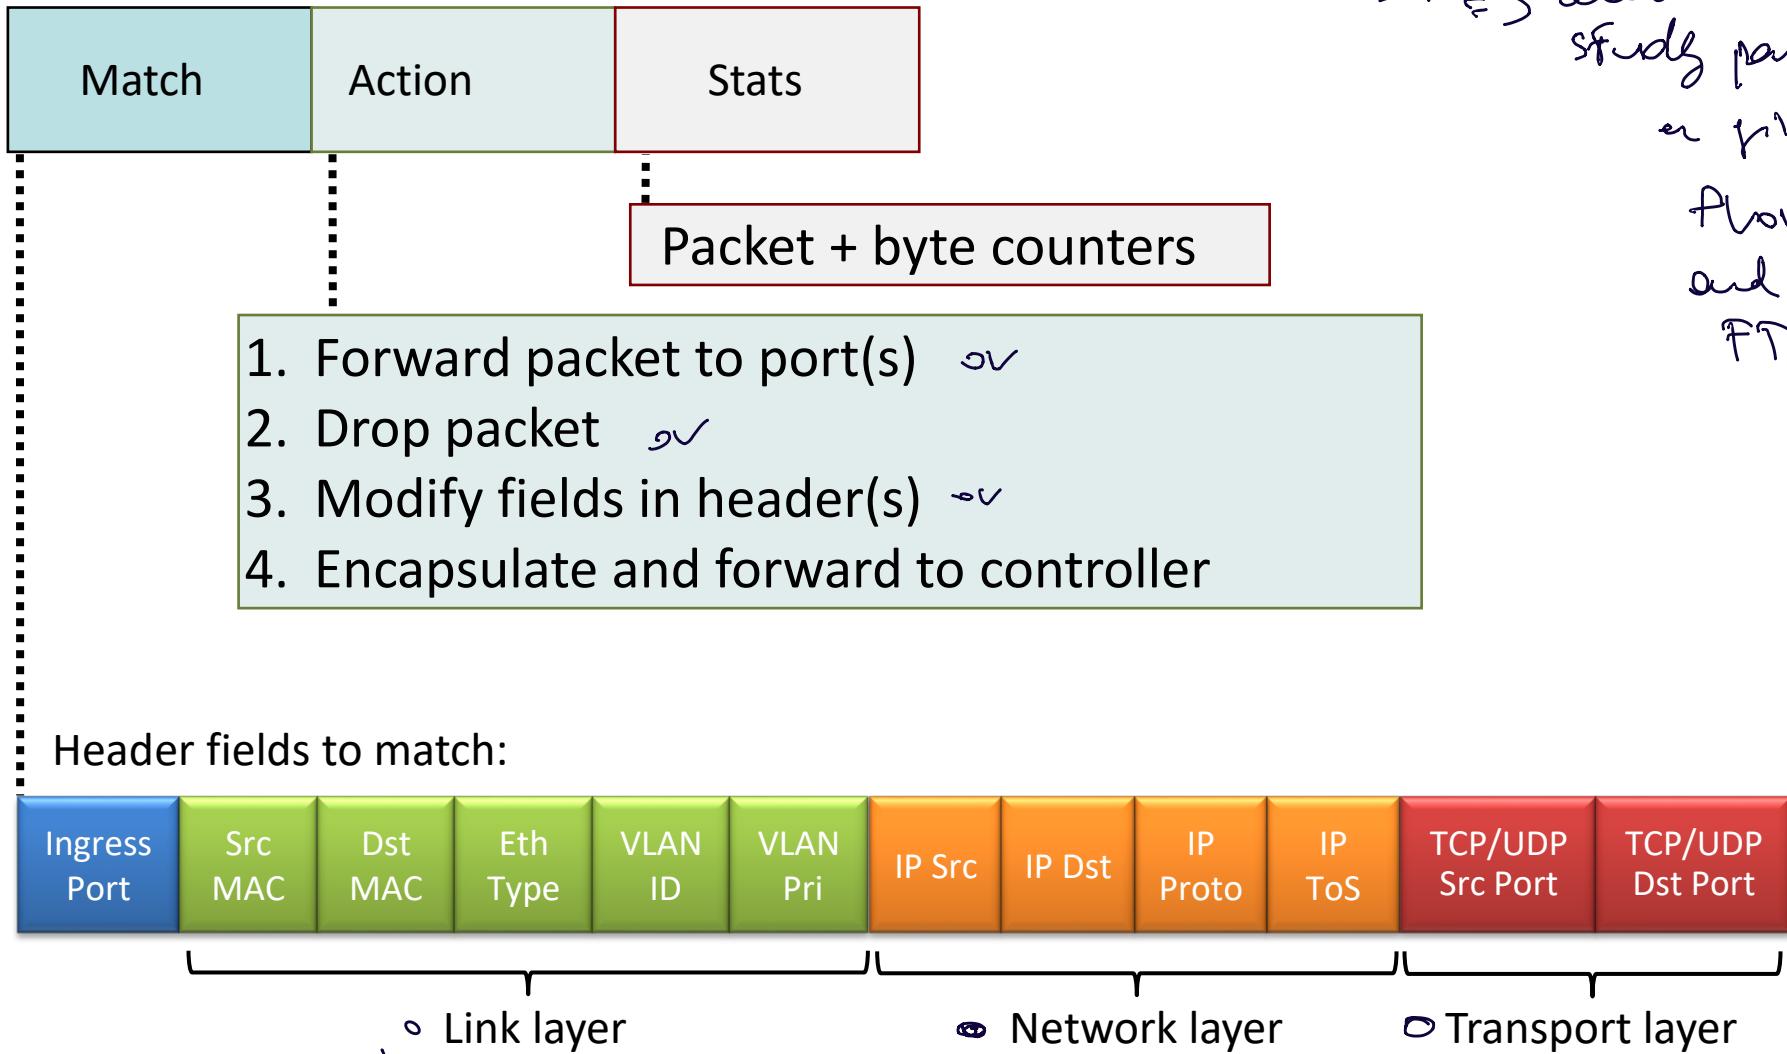


in openflow
packet IN
must
match IP source/dst

generable form,

OpenFlow: flow table entries

Packet to controller, it decides,
send back to switch the
instructions/actions to perform



~~way~~ is generalized forward

fields in the header

OpenFlow: examples

Destination-based forwarding: but we specific the dest address to forward at port: X

Switch Port	MAC src	MAC dst	Eth type	VLAN ID	VLAN Pri	IP Src	IP Dst	IP Prot	IP ToS	TCP s-port	TCP d-port	Action
*	*	*	*	*	*	*	51.6.0.8	*	*	*	*	port6

IP datagrams destined to IP address 51.6.0.8 should be forwarded to router output port 6

Firewall:

Switch Port	MAC src	MAC dst	Eth type	VLAN ID	VLAN Pri	IP Src	IP Dst	IP Prot	IP ToS	TCP s-port	TCP d-port	Action
*	*	*	*	*	*	*	*	*	*	*	22	drop

Block (do not forward) all datagrams destined to TCP port 22 (ssh port #)

Switch Port	MAC src	MAC dst	Eth type	VLAN ID	VLAN Pri	IP Src	IP Dst	IP Prot	IP ToS	TCP s-port	TCP d-port	Action
*	*	*	*	*	*	128.119.1.1	*	*	*	*	*	drop

Block (do not forward) all datagrams sent by host 128.119.1.1 (*suspicious host*)

OpenFlow: examples

behavior of a switch

Layer 2 destination-based forwarding:

Switch Port	MAC src	MAC dst	Eth type	VLAN ID	VLAN Pri	IP Src	IP Dst	IP Prot	IP ToS	TCP s-port	TCP d-port	Action
-------------	---------	---------	----------	---------	----------	--------	--------	---------	--------	------------	------------	--------

* * 22:A7:23:
11:E1:02 * * * * * * * * * * * port3

layer 2 frames with destination MAC address 22:A7:23:11:E1:02 should be forwarded to output port 3

OpenFlow abstraction behaviors implemented:

match+action: abstraction unifies different kinds of devices

Router

- *match:* longest destination IP prefix
- *action:* forward out a link

Switch

- *match:* destination MAC address
- *action:* forward or flood
To port *more ports*

Firewall

- *match:* IP addresses and TCP/UDP port numbers
- *action:* permit or deny (allow/drop)

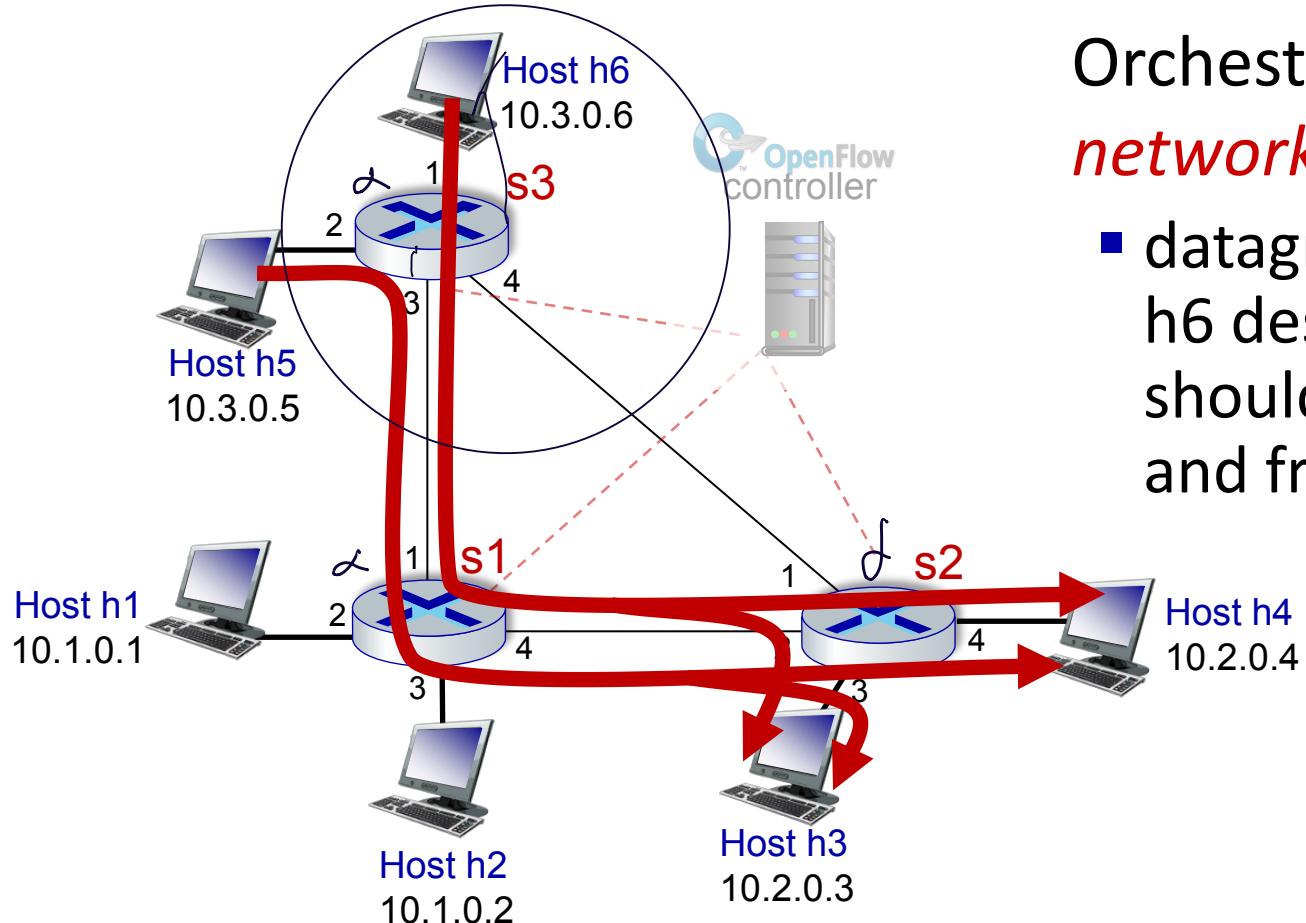
NAT

- *match:* IP address and port
- *action:* rewrite address and port

device can modify the header:

- 1st write a new ip port (mystic)
- 2nd rewrite/overwrite ip address

OpenFlow example

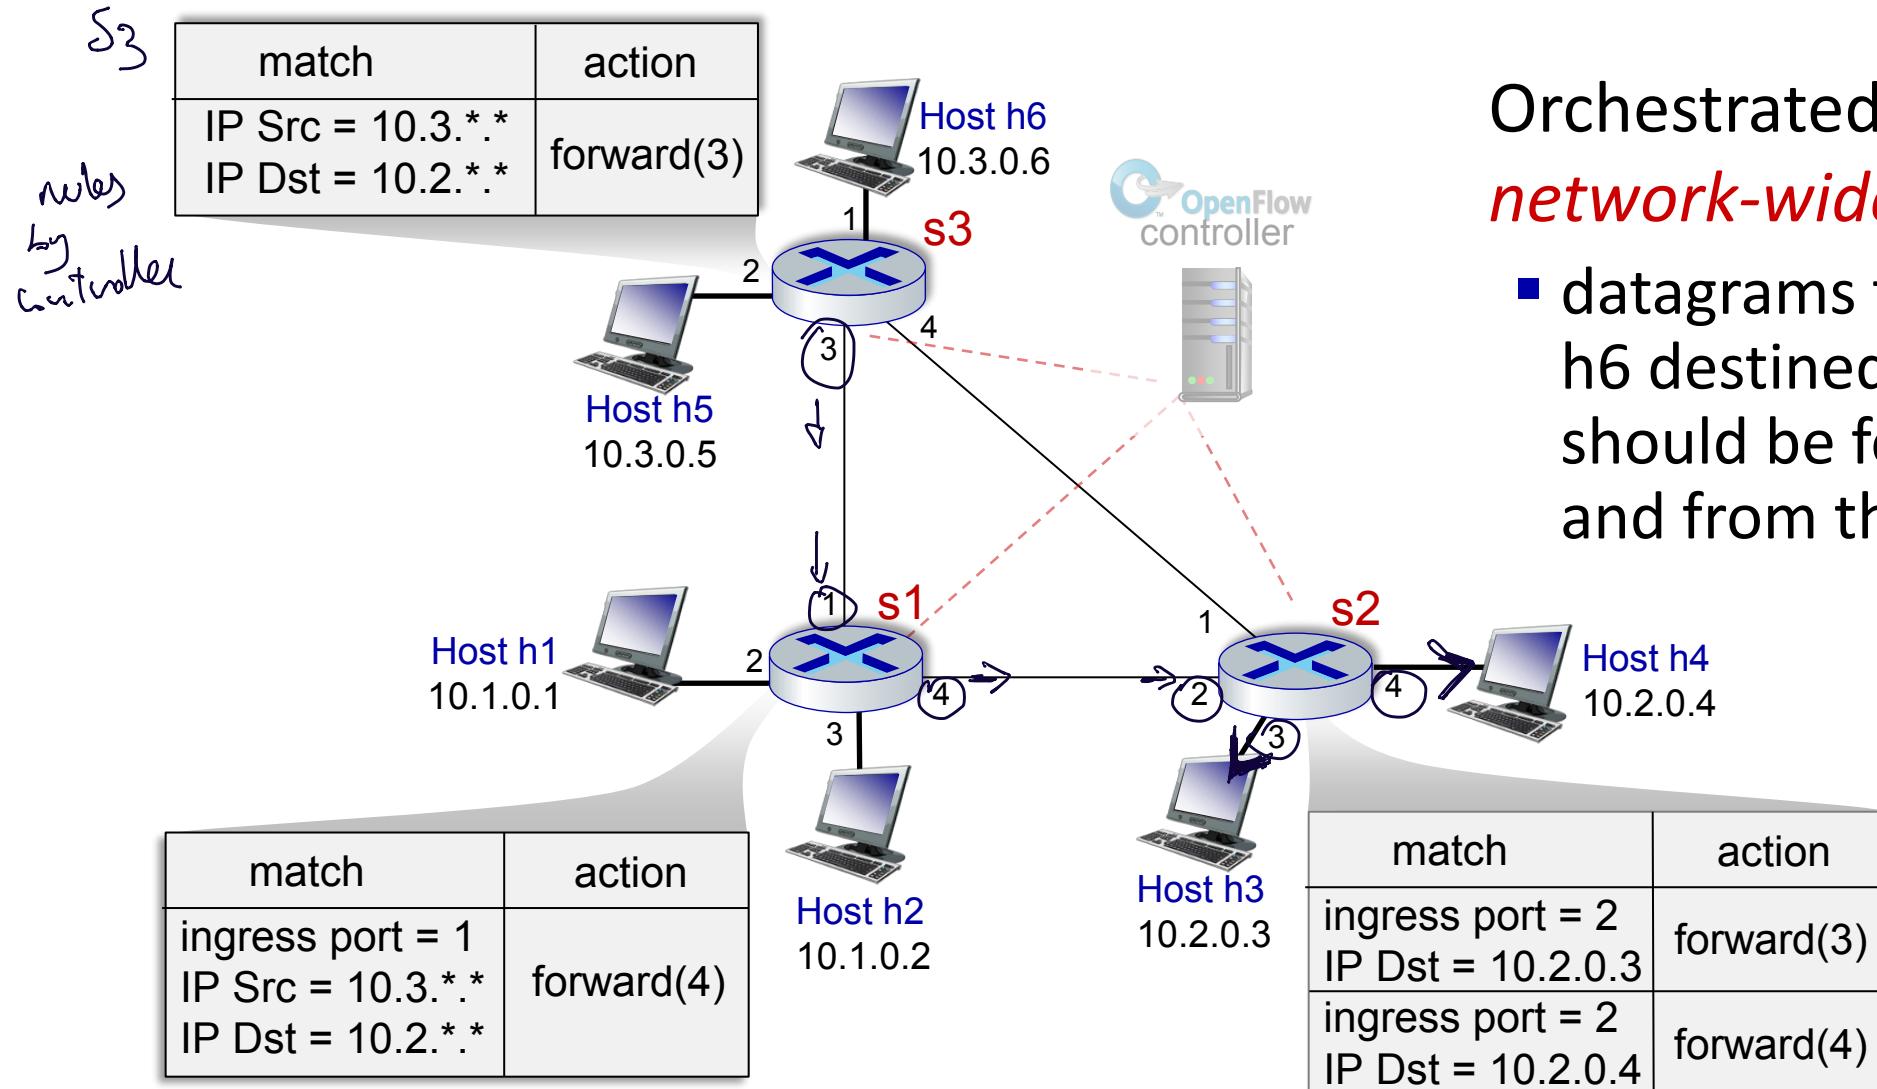


Orchestrated tables can create **network-wide** behavior, e.g.:

- datagrams from hosts h5 and h6 destined to h3 or h4, should be forwarded via s1 and from there to s2

of course rules:
from S1/S2
to S3/S4
rule: PATT&R

OpenFlow example



Orchestrated tables can create **network-wide** behavior, e.g.:

- datagrams from hosts h5 and h6 destined to h3 or h4, should be forwarded via s1 and from there to s2

→ controller:
orchestrate
rules to
implement
behavior on
the complete
topology of the
network

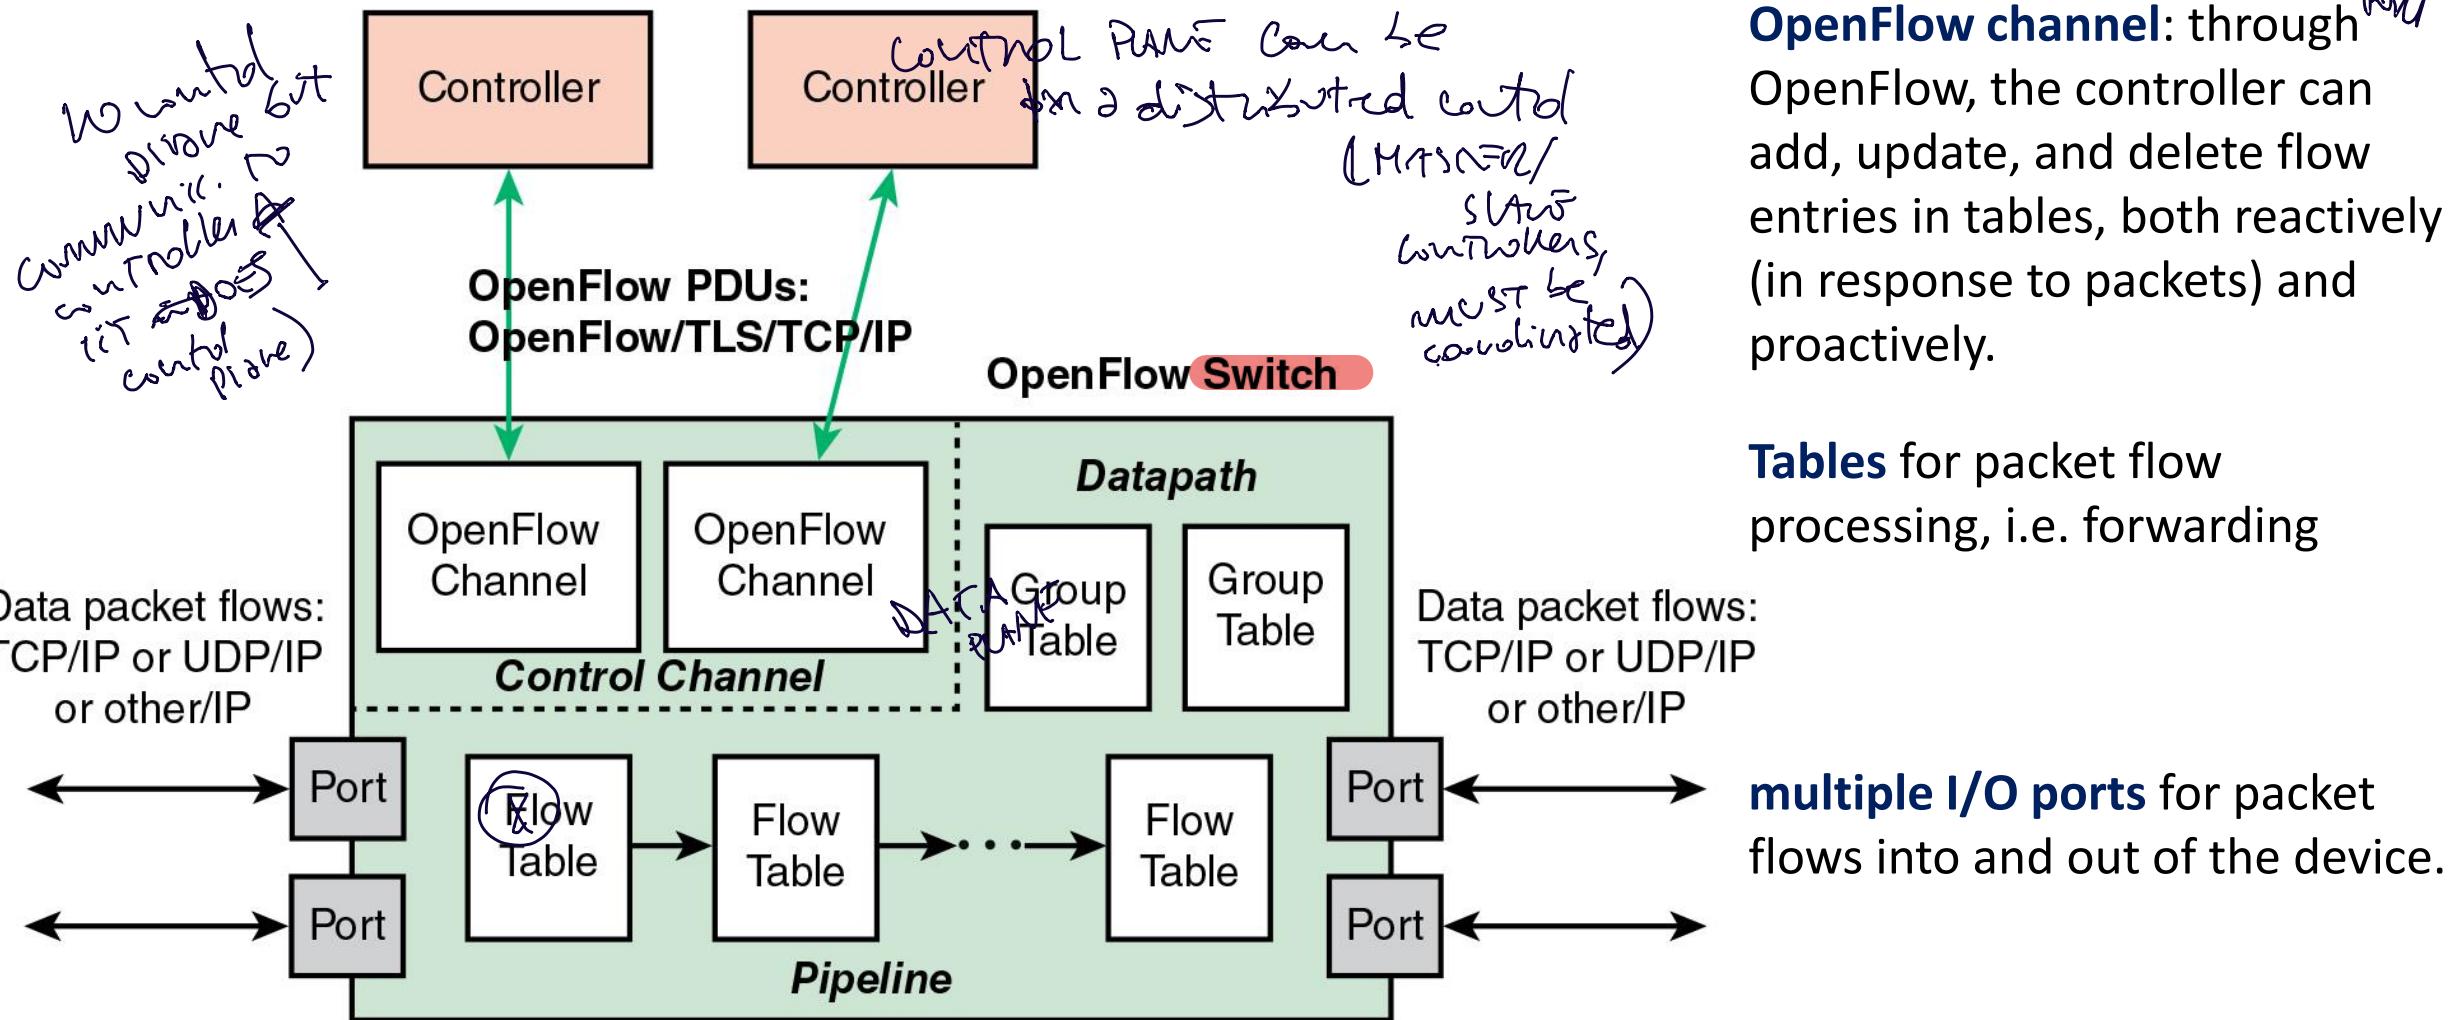
Generalized forwarding: summary

in flow fields

- “match plus action” abstraction: match bits in arriving packet header(s) in any layers, take action
 - matching over many fields (link-, network-, transport-layer)
 - local actions: drop, forward, modify, or send matched packet to controller
 - “program” network-wide behaviors
- simple form of “network programmability” *(instead of AP), based on FT*
 - programmable, per-packet “processing”
 - *historical roots: active networking*
 - *today: more generalized programming:*
P4 (see p4.org) *program to execute to implement specific behaviors*

OpenFlow Switch

- group table: specify entries of flow
1/m flows of
2 groups
1
0pt
Port



OpenFlow channel: through OpenFlow, the controller can add, update, and delete flow entries in tables, both reactively (in response to packets) and proactively.

Tables for packet flow processing, i.e. forwarding

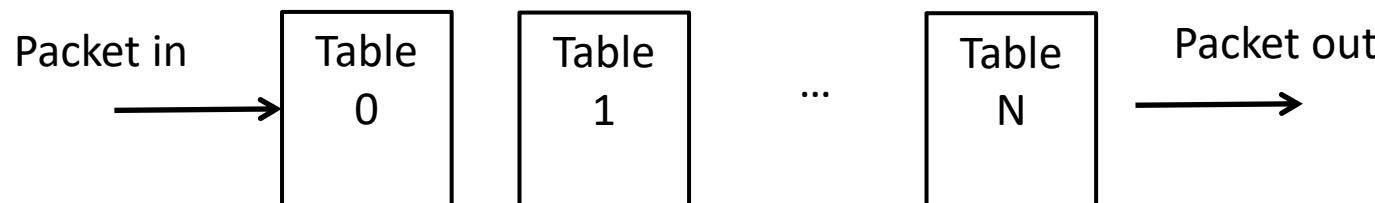
multiple I/O ports for packet flows into and out of the device.

Flow Table Pipeline

- PACKET UNIT goes authorized firstly by TAB ϕ :
 - according to ip call
 - Jump by the rules specified

- A switch includes one or more flow tables
- If there is more than one flow table, they are organized as a pipeline
 - The flow tables of an OpenFlow switch are sequentially numbered, starting at 0.
 - The packet is first matched against flow entries of flow table 0. Other flow tables may be used depending on the outcome of the match in the first table.

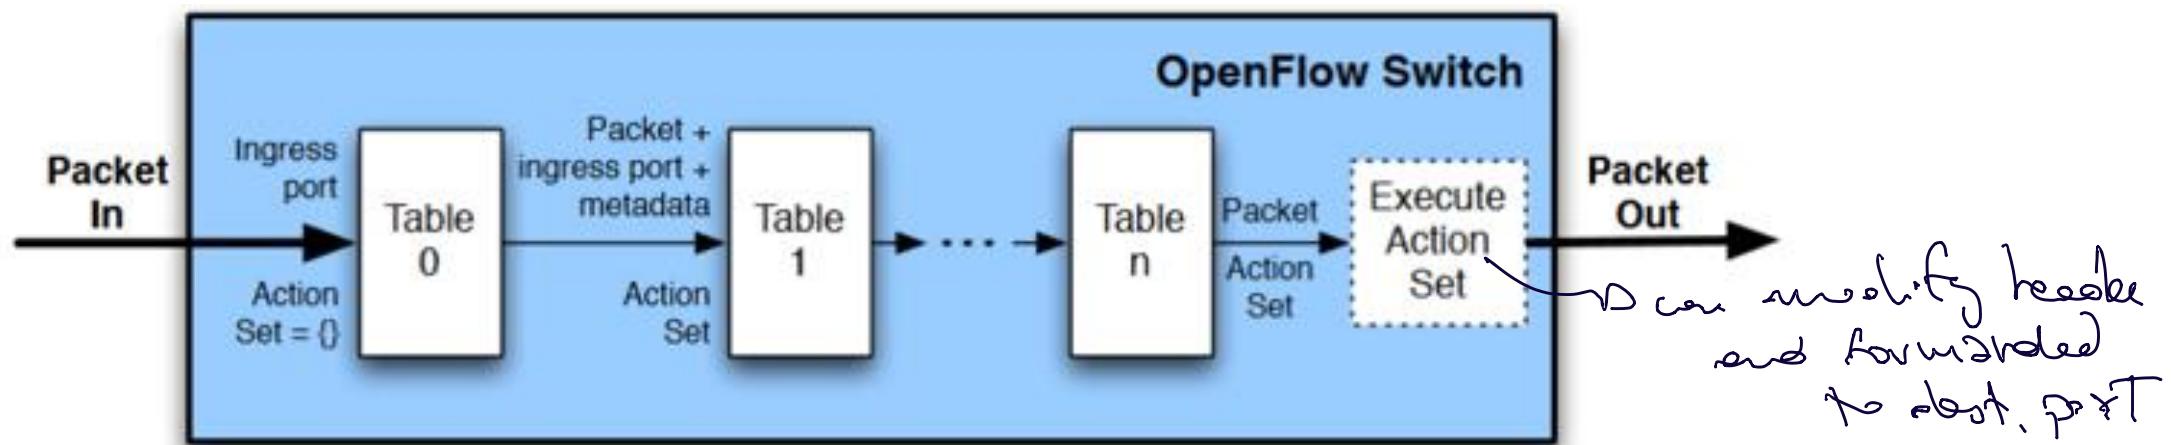
The use of multiple tables in a pipeline, rather than a single flow table, provides the SDN controller with considerable flexibility



Processing pipeline

NOT just packet goes thru table, but analyzed/processed
switch takes packet info decisions
(acting)

- When processed by a flow table, the packet is matched against the flow entries of the flow table.
- If there is a matching, any instructions associated with this entry is executed. This may include updating the action set, updating the metadata value, and performing actions.
- those instructions may explicitly direct the packet to another flow table (Goto Instruction), where the same process is repeated again, to the group table, to the meter table, or, finally, to an output port for forwarding



Processing pipeline

Problem: need to provide default actions that don't match any pattern
(shape topology, host correction, shape...)

If there is a match on one or more entries in a table the match is defined to be with the highest-priority matching entry

If there is a match only on a **table-miss entry**, the table entry may contain instructions, as with any other entry. In practice, the table-miss entry specifies one of three actions:

- a) **Send packet to controller.** This will enable the controller to define a new flow for this and similar packets, or decide to drop the packet.
- b) **Direct packet to another flow table** → specific for this
- c) **Drop the packet.**

If there is no match on any entry and there is no table-miss entry, the packet is dropped.

END of DATA PLANE

Questions

1. How does general forwarding differ from destination-based forwarding?



MATCH & ACTION

~Actions not just based on IP address but any field of header frame, can abstract behaviours of routers, switches, FIREWALL, NAT

2. What is meant by the "match-action" operation of a router or switch? In the case of destination-based forwarding, what is compared and what action is taken?

3. In the case of SDN, name three fields that can be compared and three actions that can be taken.

ACTIONS: copy, drop,
modify, log,
forward

- mac source Adr / Dest
- IP Some / Dest
- TCP port / not port

Problem

Consider the OpenFlow SDN network in the figure.

Suppose the desired forwarding behavior for datagrams arriving at s2 is as follows:

- a) datagrams arriving at port 1 from hosts h5 or h6 and having hosts h1 or h2 as destination must be forwarded to output port 2;
- b) datagrams arriving at port 2 from hosts h1 or h2 and having hosts h5 or h6 as destination must be forwarded to output port 1;
- c) datagrams arriving at port 1 or 2 and having hosts h3 or h4 as destination must be forwarded to the specified host;
- d) hosts h3 or h4 must be able to exchange datagrams.

Specify occurrences of the s2 flow table to implement this behavior

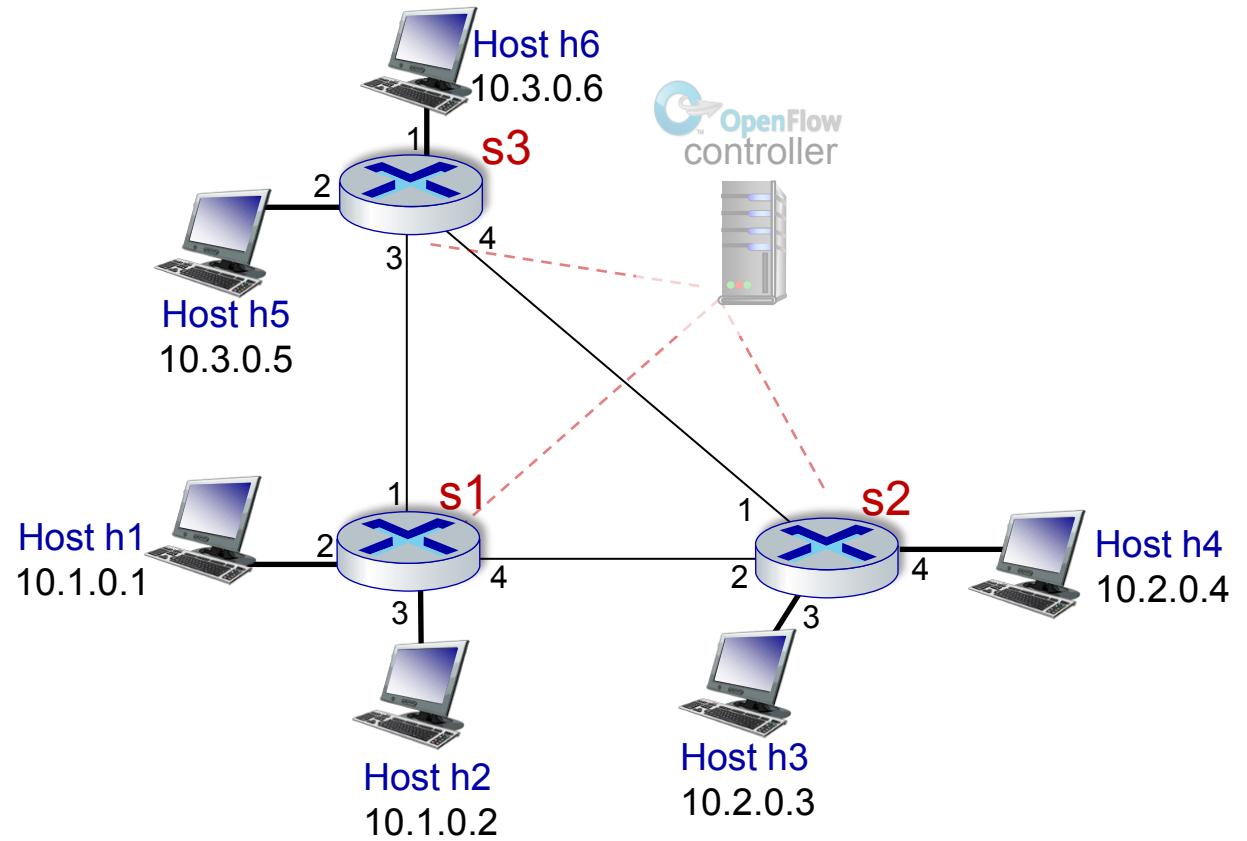
MATCH

iPSNC = 10.3.*.* , inner port = 1,

ACTION

FORWARD(2)

IP DEST: 10.1.0.*



References

Prof. Scott Shenker of UC Berkeley, [*"The Future of Networking, and the Past of Protocols"*](#)

Stallings, W. (2015). Foundations of modern networking: SDN, NFV, QoE, IoT, and Cloud. Addison-Wesley Professional.

Shenker, S., Casado, M., Koponen, T., & McKeown, N. (2011). The future of networking, and the past of protocols. Open Networking Summit, 20, 1-30.

OpenFlow switch specifications, Open Networking, v 1.5.1 2015, <https://www.opennetworking.org/software-defined-standards/specifications/>

Sushant Jain et al. 2013. B4: experience with a globally-deployed software defined wan. In Proceedings of the ACM SIGCOMM 2013 conference on SIGCOMM (SIGCOMM '13). ACM, New York, NY, USA, 3-14

A Service-Oriented Deployment Policy of End-to-End Network Slicing Based on Complex Network Theory - Scientific Figure on ResearchGate. Available from: https://www.researchgate.net/figure/5G-network-slicing-architecture_fig1_324175599 [accessed 13 Mar, 2019]

<http://thenetworksherpa.com/data-control-plane-separation-sortof/>

Paul Goransson, Chuck Black, Software Defined Networks: A Comprehensive Approach, Morgan Kaufmann

Marco Cello, Talk@IEIIT CNR, Genova, 28 Marzo 2014

Sushant Jain, Alok Kumar, Subhasree Mandal, Joon Ong, Leon Poutievski, Arjun Singh, Subbaiah Venkata, Jim Wanderer, Junlan Zhou, Min Zhu, Jon Zolla, Urs Hözle, Stephen Stuart, and Amin Vahdat, B4: experience with a globally-deployed software defined wan. In Proceedings of the ACM SIGCOMM 2013 conference on SIGCOMM (SIGCOMM '13). ACM, New York, NY, USA, 3-14

From: SDN Apps (or “how I can use SDN in my daily life”) Dr. Vasileios Kotronis vkotronis@ics.forth.gr