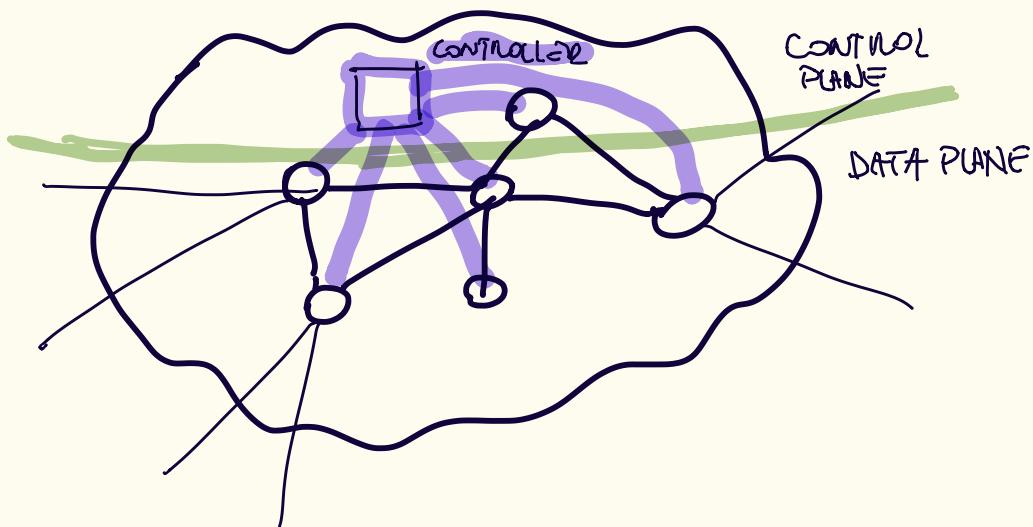


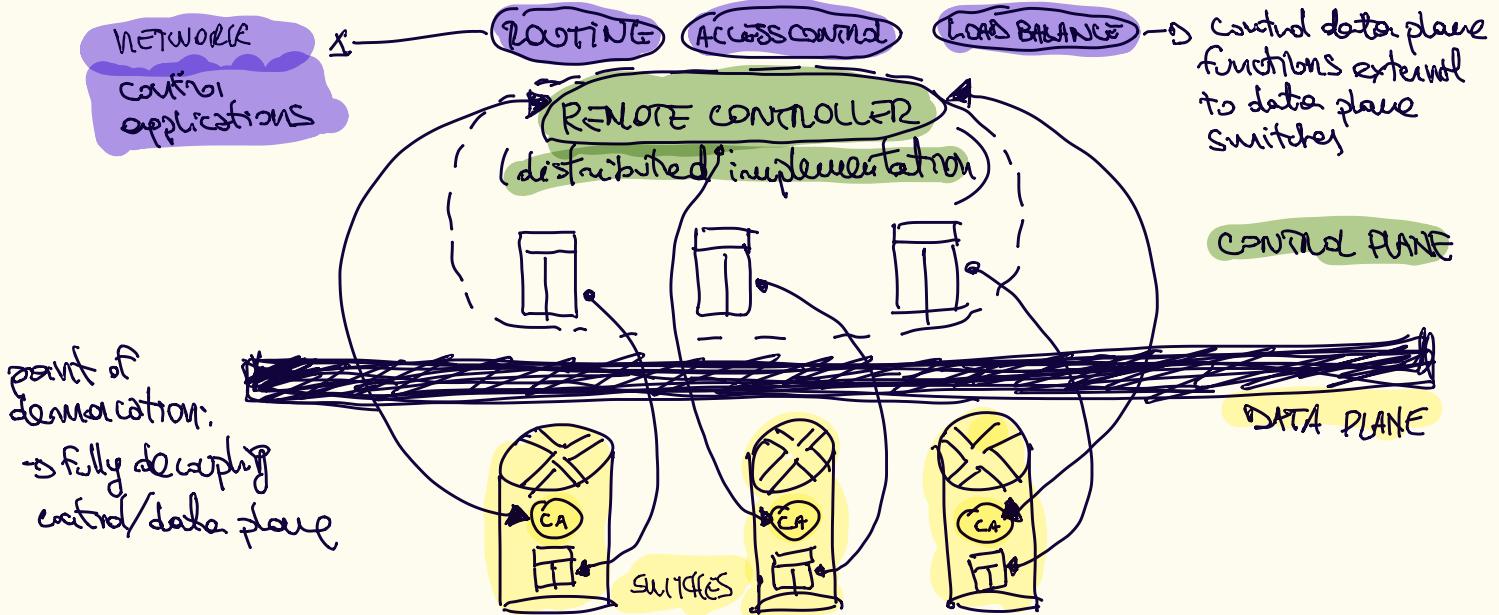
SDN: Software defined networking



Data plane: forwarding traffic

Control plane: compute state of routers that allows to make decisions and forward traffic to the destination: state \Rightarrow routing tables
Network control plane computes the routing table

- In today's networking: Data & Control planes run on routers \rightarrow distributed in the network
- In SDN: CONTROL PLANE runs in a **LOGICALLY CENTRALIZED CONTROLLER**
 - \Rightarrow separation of control plane & data plane
 - \Rightarrow allows to build a network with commodity devices where control resides in a separate control program
 - \Rightarrow controls routers across the network
 - \Rightarrow control program exerts control over the routers in the network
 - \hookrightarrow facilitates network-wide control
 - \Rightarrow this refactoring allows to move from a network where devices are vertically integrated to a network where devices have open interfaces that can be controlled by SW
- **Data plane** \Rightarrow forwarding; moving packets from router's input port to routers' output port, according to F.
- **Control plane** \Rightarrow **compute FT**: determine route taken by packet from source to destination
 - classic: per-router control plane: individual routing algo in each router interact in control plane to compute routing tables.
 - SDN: logically centralized control plane: remote controller computes/install FT in routers



- Remote controllers need replicas to manage false positives
- Dedicated devices (servers) to compute FT
- Receive a packet, look into FT, deliver to destination
- Remote controller collects infos to get network state, sending instructions to routers
- Routers talk to RC, infos on paths/links/statistics

□ Why a logically CONTROL PLANE?

CONTROL PLANE

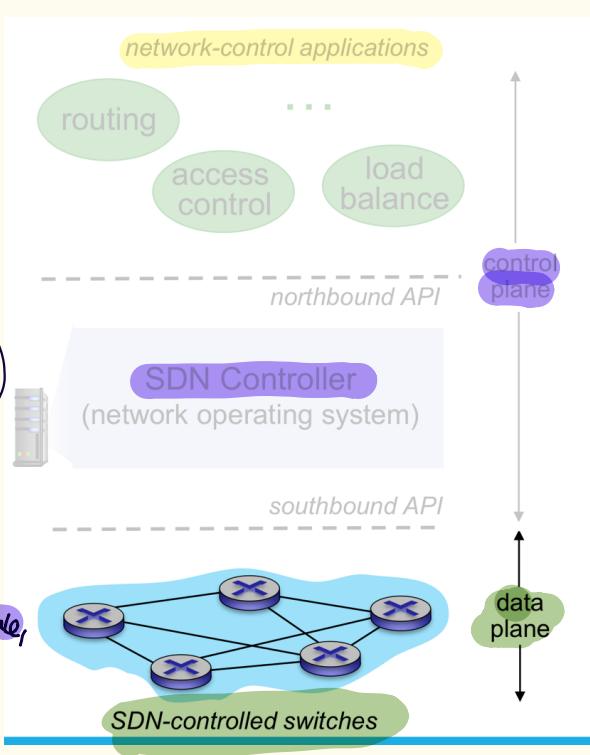
- 1) Easier NETWORK management, traffic flows flexibility
- 2) More easy defining applications to handle the traffic
- 3) Possibility to set different routing strategies
- 4) (CENTRAL CONCEPT): table-based forwarding (OPENFLOW API) allows PROGRAMMING FORWARDING
 - centralized programming easier: compute tables centrally and distributes
 - distributed programming harder: compute tables as result of distributed algo implemented in each and every router

1) SDN controlled switches: DATA PLANE

- 1) Switches implementing data-plane forwarding in hardware
- 2) Flow table computed, installed by controller
- 3) API for table-based switch control (OpenFlow)
 - receive FT by controller and exploit API so table can be updated
- 4) Protocol communicating with controller (SOUTH API)

2) SDN CONTROLLER (NETWORK OS)

- 1) Builds topology on underlying level
- 2) Decides routes (FT)
- 3) Maintain network STATE information
- 4) Implemented as distributed system for performance, scalability, robustness, fault-tolerance
- 5) Interacts with NETWORK-CONTROL APPS via northbound API and with SDN-CONTROLLED SWITCHES via southbound API



3) NETWORK-CONTROL ABSTRACTION

1) Branches of control:

- implements control functions using lower-level services (API provided by controller)
- northbound API: high level of abstraction than southbound API

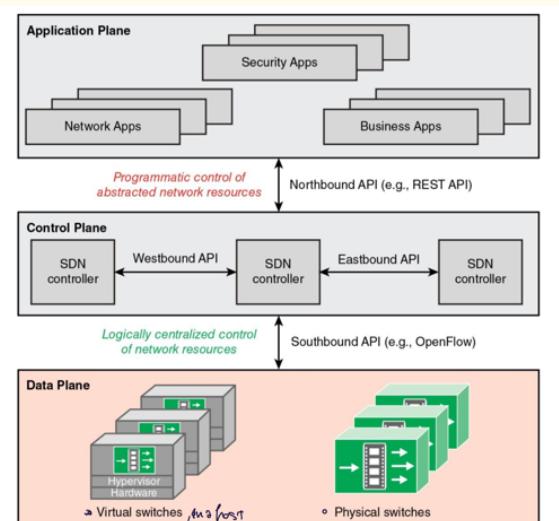
2) "UNBOUNDED": can be provided by 3rd parties: distinct from vendor vendors or SDN controller

DATA PLANE

• FORWARDING ABSTRACTION

- Resource or infrastructure layer made by forwarding devices (forwarding switches)
 - perform transport and processing of data according to decisions made by the SDN control plane
 - without embedded SW to make autonomous decisions

• FUNCTIONS



1) DATA FORWARDING: look in the FT to match the traffic flow along data forwarder paths, computed and established according to rules by SDN apps.

2) CONTROL SUPPORT: interactions with SDN controller for managing forwarding rules

- OpenFlow switch control: module to interact with controller

GENERALIZED FORWARDING: MATCH + ACTIONS

"match + action" abstraction: performed on packets, match bits in source packets, take actions:

1) DESTINATION-BASED FORWARDING: forward based on IP address

2) GENERALIZED FORWARDING:

- many header fields can determine actions (not necessarily an IP/Port match, other fields, like MAC address)
- More actions possible: DROP/COPY/MODIFY/LOG (different actions in FT)

packets

FLOW TABLE ABSTRACTION

1) OFLOW: defined by header fields values

- header's FT defines router's match + action rules

Fields:



src = *.*.*.* dest=3.4.*.*
src=1.2.*.* dest=*.**.*
src=10.1.2.3 dest=*.**.*

- should keep FT short as possible

forward(2)
drop
send to controller

- if you don't have a specific behaviour, switches send packet to controller, it will decide actions on that, sending back to the switch instructions to perform

2) GENERALIZED FORWARDING: simple packet handling rules

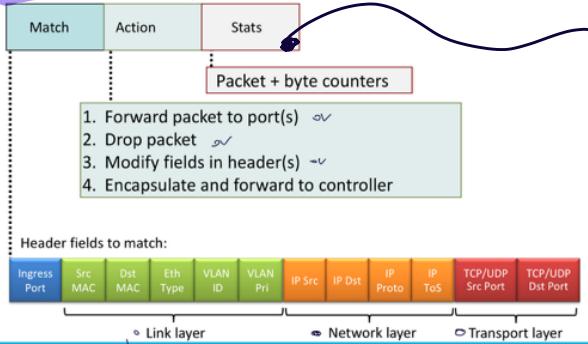
1) MATCH: pattern values in packet header fields

2) ACTIONS: for matched packets: drop, forward, modify matched packets or send matched packets to controller that tells actions to be performed on that packet

3) PRIORITY: disambiguate overlapping patterns: can have more matches, with higher priority

4) COUNTERS: #bytes and # of packets for statistics

OPENFLOW: FLOW TABLE ENTRIES



Destination-based forwarding:

Switch Port	MAC src	MAC dst	Eth type	VLAN ID	VLAN Pri	IP Src	IP Dst	IP Prot	IP ToS	TCP s-port	TCP d-port	Action
*	*	*	*	*	*	*	51.6.0.8	*	*	*	*	port6

IP datagrams destined to IP address 51.6.0.8 should be forwarded to router output port 6

• Can specify dest address to forward at port x

• BEHAVIOR OF A SWITCH:

Layer 2 destination-based forwarding:

Switch Port	MAC src	MAC dst	Eth type	VLAN ID	VLAN Pri	IP Src	IP Dst	IP Prot	IP ToS	TCP s-port	TCP d-port	Action
*	*	22:A7:23:11:E1:02	*	*	*	*	*	*	*	*	*	port3

layer 2 frames with destination MAC address 22:A7:23:11:E1:02 should be forwarded to output port 3

OPENFLOW ABSTRACTIONS: the implemented behaviors

Router

- **match:** longest destination IP prefix
- **action:** forward out a link

Switch

- **match:** destination MAC address
- **action:** forward or flood

Firewall

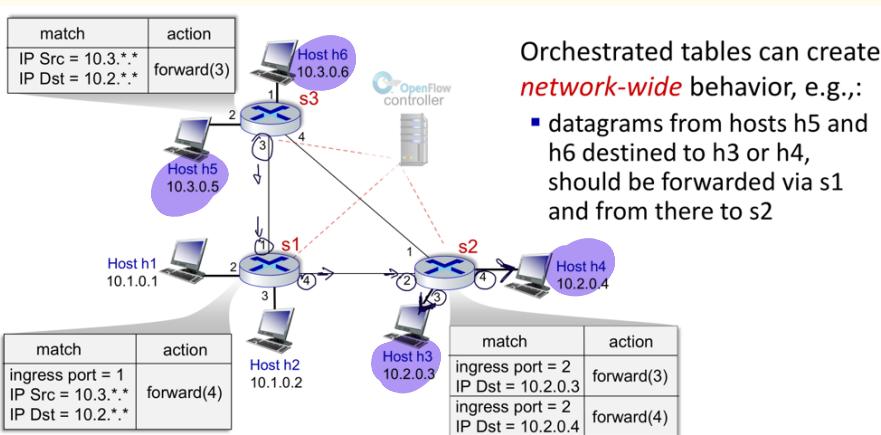
- **match:** IP addresses and TCP/UDP port numbers
- **action:** permit or deny

NAT

- **match:** IP address and port
- **action:** rewrite address and port

→ devices can modify the header
 - 1st: match to a new private ip/port
 - 2nd: rewrite/substitute ip address with public one

OPENFLOW EXAMPLE



Orchestrated tables can create **network-wide** behavior, e.g.:

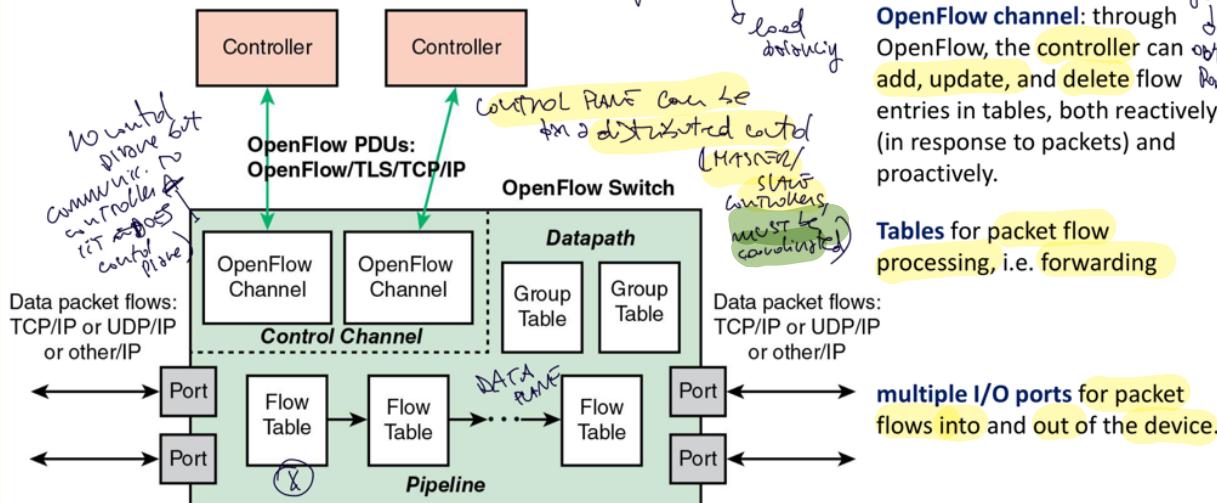
- datagrams from hosts h5 and h6 destined to h3 or h4, should be forwarded via s1 and from there to s2

- Controller orchestrates rules to implement the behavior of the network topology

SUMMARY ON GENERALIZED FORWARDING

- “match plus action” abstraction: match bits in arriving packet header(s) in any layers, take action
 - matching over many fields (link-, network-, transport-layer)
 - local actions: drop, forward, modify, or send matched packet to controller
 - “program” network-wide behaviors
 - simple form of “network programmability” install the AP, control via FT
 - programmable, per-packet “processing”
 - historical roots: active networking
 - today: more generalized programming:
- P4 (see p4.org): program to create to implement specific behaviors!

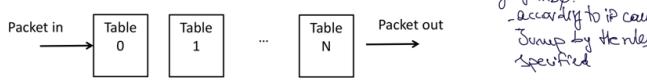
OpenFlow Switch



Flow Table Pipeline

- A switch includes one or more flow tables
- If there is more than one flow table, they are organized as a pipeline
- The flow tables of an OpenFlow switch are sequentially numbered, starting at 0.
- The packet is first matched against flow entries of flow table 0. Other flow tables may be used depending on the outcome of the match in the first table.

The use of multiple tables in a pipeline, rather than a single flow table, provides the SDN controller with considerable flexibility



2)

Processing pipeline

If there is a match on one or more entries in a table the match is defined to be with the **highest-priority** matching entry

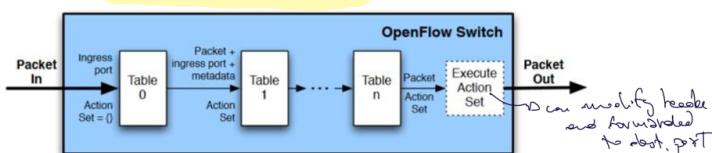
If there is a match only on a table-miss entry, the table entry may contain instructions, as with any other entry. In practice, the table-miss entry specifies one of three actions:

- Send packet to controller. This will enable the controller to define a new flow for this and similar packets, or decide to drop the packet.
- Direct packet to another flow table further down the pipeline.
- Drop the packet.

If there is no match on any entry and there is no table-miss entry, the packet is dropped.

Processing pipeline

- When processed by a flow table, the packet is matched against the flow entries of the flow table.
- If there is a matching, any instructions associated with this entry is executed. This may include updating the action set, updating the metadata value, and performing actions.
- those instructions may explicitly direct the packet to another flow table (Goto Instruction), where the same process is repeated again, to the group table, to the meter table, or, finally, to an output port for forwarding



Problem: need to provide default actions that don't match any pattern
 (shape topology, host correction, drop...)
 (best controller)
 => STAB miss entry (wildcard match)
 3 options / actions

END of DATA PLANE